

< Deep Learning - PART1 TF2 Basics >

Ch 1. Workshop - TensorFlow 2.0 Installation & Testing

2021/10/01

<< Installation of TF 2.0 with Anaconda3 >>

- First, install Anaconda 3 for Windows/macOS/Linux from <https://www.anaconda.com/distribution/> (<https://www.anaconda.com/distribution/>)
- Next, run TensorFlow 2 (for CPU) Setup on Anaconda Prompt :

```
conda install tensorflow
```

[Reference]:

- TensorFlow.org, "Install TensorFlow 2" <https://www.tensorflow.org/install> (<https://www.tensorflow.org/install>)
- 海萨, "Anaconda 安装tensorflow 2.0 报错解决办法" <https://zhuanlan.zhihu.com/p/62031082> (<https://zhuanlan.zhihu.com/p/62031082>)
- TensorFlow.org, "Get Started with TensorFlow" <https://www.tensorflow.org/tutorials/#get-started-with-tensorflow> (<https://www.tensorflow.org/tutorials/#get-started-with-tensorflow>)

[Content]

- [1. Testing TF 2.0](#)
- [2. How to run TensorFlow 1.x code on TF 2.0](#)

1. Testing TF 2.0

In [1]:



```
1 import tensorflow as tf
2 print(tf.__version__)
```

2.4.1

In [2]:



```
1 # -----
2 # The following code is adopted from
3 # Tutorial document of TensorFlow.org
4 # for testing TensorFlow 2.0 setup:
5 #
6 # "Get Started with TensorFlow"
7 # https://www.tensorflow.org/tutorials/#get-started-with-tensorflow
8 # -----
9
10 mnist = tf.keras.datasets.mnist
11
12 (x_train, y_train), (x_test, y_test) = mnist.load_data()
13 x_train, x_test = x_train / 255.0, x_test / 255.0
14
15 model = tf.keras.models.Sequential([
16     tf.keras.layers.Flatten(input_shape=(28, 28)),
17     tf.keras.layers.Dense(512, activation=tf.nn.relu),
18     tf.keras.layers.Dropout(0.2),
19     tf.keras.layers.Dense(10, activation=tf.nn.softmax)
20 ])
21 model.compile(optimizer='adam',
22               loss='sparse_categorical_crossentropy',
23               metrics=['accuracy'])
24
25 model.fit(x_train, y_train, epochs=5)
26 model.evaluate(x_test, y_test, verbose=0) # verbose: Verbosity mode. 0=silent
```

```
Epoch 1/5
1875/1875 [=====] - 18s 9ms/step - los
s: 0.3670 - accuracy: 0.8910
Epoch 2/5
1875/1875 [=====] - 15s 8ms/step - los
s: 0.0985 - accuracy: 0.9696
Epoch 3/5
1875/1875 [=====] - 15s 8ms/step - los
s: 0.0697 - accuracy: 0.9785
Epoch 4/5
1875/1875 [=====] - 15s 8ms/step - los
s: 0.0506 - accuracy: 0.9837
Epoch 5/5
1875/1875 [=====] - 14s 8ms/step - los
s: 0.0406 - accuracy: 0.9865
```

Out[2]:

```
[0.07084144651889801, 0.9793000221252441]
```

2. How to run TensorFlow 1.x code on TF 2.0

- It is still possible to run 1.X code, unmodified (except for contrib), in TensorFlow 2.0:

```
import tensorflow.compat.v1 as tf

tf.disable_v2_behavior()
```

[NOTE]:

- More detailed information regarding " Migrate your TensorFlow 1 code to TensorFlow 2 " can be found here: <https://www.tensorflow.org/guide/migrate> (<https://www.tensorflow.org/guide/migrate>).

In [7]:



```
1 import tensorflow.compat.v1 as tf
2 tf.disable_v2_behavior()
3
4 print(tf.__version__)
```

WARNING:tensorflow:From C:\Users\USER\Anaconda3\lib\site-packages\tensorflow_core\python\compat\v2_compat.py:65: disable_resource_variables (from tensorflow.python.ops.variable_scope) is deprecated and will be removed in a future version.
Instructions for updating:
non-resource variables are not supported in the long term
2.0.0

The following code is adopted for testing TensorFlow 2.0 setup from the reference below:

- Tom Hope, Yehezkel S. Resheff, and Itay Lieder, "**Learning TensorFlow : A Guide to Building Deep Learning Systems**," Chapter 2 & 4, O'Reilly, 2017. <https://goo.gl/iEmehh> (<https://goo.gl/iEmehh>)
 - Download the code from GitHub : <https://github.com/gigwegbe/Learning-TensorFlow> (<https://github.com/gigwegbe/Learning-TensorFlow>)
-

Loading the MNIST dataset (from TensorFlow 2.0)

In [8]:



```
1 mnist = tf.keras.datasets.mnist
2
3 (x_train, y_train), (x_test, y_test) = mnist.load_data()
4 x_train, x_test = x_train / 255.0, x_test / 255.0
```

In [9]:



```
1 import numpy as np
2
3 x_train = np.array([x_train[i].flatten() for i in range(len(x_train))])
4 x_train.shape
```

Out[9]:

(60000, 784)

In [10]:



```
1 x_test = np.array([x_test[i].flatten() for i in range(len(x_test))])
2 x_test.shape
```

Out[10]:

(10000, 784)

In [11]:



```
1 y_train[0], y_test[0]
```

Out[11]:

(5, 7)

In [12]:



```
1 def one_hot(vec, vals=10):
2     n = len(vec)
3     out = np.zeros((n, vals))
4     out[range(n), vec] = 1
5     return out
```

In [13]:



```
1 y_train = one_hot(y_train)
2 y_train[0]
```

Out[13]:

array([0., 0., 0., 0., 0., 1., 0., 0., 0., 0.])

In [14]:

```
1 y_test = one_hot(y_test)
2 y_test[0]
```

Out[14]:

```
array([0., 0., 0., 0., 0., 0., 0., 1., 0., 0.])
```

Building a Computation Graph on TF 1.x

In [15]:

```
1 # Each Input Image, X, with 28*28 (= 784) pixels
2 X = tf.placeholder(tf.float32, [None, 784])
3
4 # y_true : the training Labeled dataset
5 y_true = tf.placeholder(tf.float32, [None, 10])
```

In [16]:

```
1 # Initializing Weights & Biases for Nodes in ALL Hidden Layers
2 def weight_variable(shape):
3     initial = tf.truncated_normal(shape, stddev=0.1)
4     return tf.Variable(initial)
5
6 def bias_variable(shape):
7     initial = tf.constant(0.1, shape=shape)
8     return tf.Variable(initial)
```

In [17]:

```
1 # Building a Fully-Connected Deep Network
2 def full_layer(inputs, size):
3     in_size = int(inputs.get_shape()[1])
4     W = weight_variable([in_size, size])
5     b = bias_variable([size])
6     return tf.add(tf.matmul(inputs, W), b)
```

In [18]:



```
1 keep_prob = tf.placeholder(tf.float32)
2
3 # < Hidden Layer 1 >
4 layer_1_drop = tf.nn.dropout(X, keep_prob=keep_prob)
5 # Activation Function : ReLU
6 layer_1_Outputs = tf.nn.relu(full_layer(layer_1_drop, 256))
7
8 # < Hidden Layer 2 >
9 layer_2_drop = tf.nn.dropout(layer_1_Outputs, keep_prob=keep_prob)
10 # Activation Function : ReLU
11 layer_2_Outputs = tf.nn.relu(full_layer(layer_2_drop, 128))
12
13 # < Output Layer >
14 output_drop = tf.nn.dropout(layer_2_Outputs, keep_prob=keep_prob)
15 # Without Activation Function
16 y_pred = full_layer(output_drop, 10)
```

WARNING:tensorflow:From <ipython-input-18-931f684597d4>:4: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

In [19]:



```
1 cross_entropy = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=
2 gd_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
3
4 correct_mask = tf.equal(tf.argmax(y_pred, 1), tf.argmax(y_true, 1))
5 accuracy = tf.reduce_mean(tf.cast(correct_mask, tf.float32))
```

WARNING:tensorflow:From <ipython-input-19-315e39f82a5d>:1: softmax_cross_entropy_with_logits (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:

Future major versions of TensorFlow will allow gradients to flow into the labels input on backprop by default.

See `tf.nn.softmax_cross_entropy_with_logits_v2`.

Launching the Computation Graph on TF 1.x

In [20]:



```
1 def next_batch(i, images, labels, batch_size):
2     i_start = (i * batch_size) % len(images)
3     x, y = images[i_start : i_start+batch_size], labels[i_start : i_start+batch_size]
4     return x, y
```

In [21]:



```
1 NUM_STEPS = 8000
2 MINIBATCH_SIZE = 100
3 Display_Step = 1000
4
5 with tf.Session() as sess:
6     sess.run(tf.global_variables_initializer())
7
8     for i in range(NUM_STEPS):
9         batch_xs, batch_ys = next_batch(i, x_train, y_train, MINIBATCH_SIZE)
10        sess.run(gd_step, feed_dict={X: batch_xs,
11                                     y_true: batch_ys,
12                                     keep_prob: 0.5})
13
14        if (i+1) % Display_Step == 0:
15            # Calculate batch loss and accuracy
16            loss_temp, accu_temp = sess.run([cross_entropy, accuracy],
17                                           feed_dict={X: batch_xs,
18                                                     y_true: batch_ys,
19                                                     keep_prob: 1.0})
20
21            print("Step " + str(i+1).rjust(4) + \
22                  " : Loss = " + "{:.4f}".format(loss_temp) + \
23                  ", Accuracy = " + "{:.3f}".format(accu_temp))
24
25        print("\n Computing the test accuracy ... ", end = " ")
26
27        ## -----
28        ## Split the test procedure into 10 blocks of 1,000 images each.
29        ## Doing this is important mostly for much larger datasets.
30        ## -----
31        ## mnist.test.images.shape : (10000, 784)
32        X_test = x_test.reshape(10, 1000, 784)
33        ## mnist.test.labels.shape : (10000, 10)
34        Y_test = y_test.reshape(10, 1000, 10)
35
36        test_loss = np.mean([sess.run(cross_entropy,
37                                     feed_dict={X: X_test[i],
38                                               y_true: Y_test[i],
39                                               keep_prob: 1.0})
40                             for i in range(10)])
41
42        test_accu = np.mean([sess.run(accuracy,
43                                     feed_dict={X: X_test[i],
44                                               y_true: Y_test[i],
45                                               keep_prob: 1.0})
46                             for i in range(10)])
47
48        print("\n [ Test Accuracy ] : {}".format(test_accu) +
49              "\n [ Test Loss Score ] : {}".format(test_loss))
```

```
Step 1000 : Loss = 0.2206, Accuracy = 0.960
Step 2000 : Loss = 0.1248, Accuracy = 0.970
Step 3000 : Loss = 0.1113, Accuracy = 0.990
Step 4000 : Loss = 0.0721, Accuracy = 0.980
Step 5000 : Loss = 0.0949, Accuracy = 0.960
Step 6000 : Loss = 0.1253, Accuracy = 0.990
Step 7000 : Loss = 0.0690, Accuracy = 0.990
Step 8000 : Loss = 0.0743, Accuracy = 0.980
```



```
Computing the test accuracy ...  
[ Test Accuracy ] : 0.9604999423027039  
[ Test Loss Score ] : 0.12836746871471405
```