

# < Deep Learning - PART1 TF2 Basics >

## Ch 3. TensorBoard - ML Workflow Visualization

2021/10/01

[ Reference ]

1. TensorFlow Core - Tutorials : Get started with TensorBoard .

[https://www.tensorflow.org/tensorboard/get\\_started](https://www.tensorflow.org/tensorboard/get_started)  
([https://www.tensorflow.org/tensorboard/get\\_started](https://www.tensorflow.org/tensorboard/get_started))

2. Hands-on TensorBoard (TensorFlow Dev Summit 2017)

- Youtube: <https://youtu.be/eBbEDRsCmv4>  
(<https://youtu.be/eBbEDRsCmv4>)
- Code: <https://goo.gl/ZwGnPE> (<https://goo.gl/ZwGnPE>)

3. TensorBoard on Colab - TensorBoard.dev <https://tensorboard.dev/>

(<https://tensorboard.dev/>)

- [Example Code]:  
<https://colab.research.google.com/github/tensorflow/tensorboard/blob/master/doc>  
(<https://colab.research.google.com/github/tensorflow/tensorboard/blob/master/doc>)

- TensorBoard is a tool for providing the measurements and visualizations needed during the machine learning workflow.
- It enables tracking experiment metrics like loss and accuracy, visualizing the model graph, projecting embeddings to a lower dimensional space, and much more.

## < Content >

- [1. TensorBoard on localhost](#)
  - [1.1 TensorBoard with Keras Model.fit\(\)](#)
  - [1.2 TensorBoard with tf.summary](#)
- [2. Tensorboard on Google Colab](#)

## 1. TensorBoard on localhost

In [1]:



```
1 import tensorflow as tf
2 import datetime
3
4 print(tf.__version__)
```

2.4.1

## Loading the MNIST dataset

In [2]:



```
1 mnist = tf.keras.datasets.mnist
2
3 (x_train, y_train), (x_test, y_test) = mnist.load_data()
4 x_train, x_test = x_train / 255.0, x_test / 255.0
```

## Building the Keras Neural Network Model with dropout

In [3]:



```
1 model = tf.keras.models.Sequential([
2     tf.keras.layers.Flatten(input_shape=(28, 28)),
3     tf.keras.layers.Dense(512, activation='relu'),
4     tf.keras.layers.Dropout(0.2),
5     tf.keras.layers.Dense(10, activation='softmax')
6 ])
```

## 1.1 TensorBoard with Keras Model.fit()

- When training with Keras's `Model.fit()`, adding the `tf.keras.callbacks.TensorBoard` callback ensures that logs are created and stored.

- Additionally, enable histogram computation every epoch with `histogram_freq=1` (this is off by default.)

Place the logs in a timestamped subdirectory to allow easy selection of different training runs:

In [4]:



```
1 model.compile(optimizer='adam',
2               loss='sparse_categorical_crossentropy',
3               metrics=['accuracy'])
4
5 # Building Timestamped Directories For Windows :
6 log_dir="logs\\fit\\" + datetime.datetime.now().strftime('%Y%m%d_%H%M%S')
7
8 # Building Timestamped Directories For Mac/Linux :
9 # log_dir="logs/fit/" + datetime.datetime.now().strftime('%Y%m%d_%H%M%S')
10
11 # callback method for Tensorboard
12 tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
13                                                       histogram_freq=1)
14
15 model.fit(x=x_train,
16          y=y_train,
17          epochs=5,
18          validation_data=(x_test, y_test),
19          callbacks=[tensorboard_callback])
```

Epoch 1/5

1875/1875 [=====] - 23s 11ms/step - loss: 0.3651 - accuracy: 0.8932 - val\_loss: 0.1045 - val\_accuracy: 0.9686

Epoch 2/5

1875/1875 [=====] - 18s 9ms/step - loss: 0.1026 - accuracy: 0.9687 - val\_loss: 0.0810 - val\_accuracy: 0.9733

Epoch 3/5

1875/1875 [=====] - 17s 9ms/step - loss: 0.0682 - accuracy: 0.9783 - val\_loss: 0.0740 - val\_accuracy: 0.9763

Epoch 4/5

1875/1875 [=====] - 16s 9ms/step - loss: 0.0485 - accuracy: 0.9844 - val\_loss: 0.0641 - val\_accuracy: 0.9801

Epoch 5/5

1875/1875 [=====] - 17s 9ms/step - loss: 0.0409 - accuracy: 0.9862 - val\_loss: 0.0691 - val\_accuracy: 0.9788

Out[4]:

<tensorflow.python.keras.callbacks.History at 0x1dd1e103e20>

**To run TensorBoard, run the following command on Anaconda (Powershell) Prompt :**

```
tensorboard --logdir= path/to/log-directory
```

- For instance, `tensorboard --logdir logs/fit`

Connecting to `http://localhost:6006`

## 1.2 TensorBoard with `tf.summary`

- When training with methods such as `tf.GradientTape()` ([https://www.tensorflow.org/api\\_docs/python/tf/GradientTape](https://www.tensorflow.org/api_docs/python/tf/GradientTape)), use `tf.summary` to log the required information.

- Use the same dataset as above, but convert it to `tf.data.Dataset` to take advantage of batching capabilities.
- [NOTE]: The following training code is adopted from the **advanced quickstart** (<https://www.tensorflow.org/tutorials/quickstart/advanced>) tutorial.

In [5]:

```
1 train_dataset = tf.data.Dataset.from_tensor_slices((x_train, y_train))
2 test_dataset = tf.data.Dataset.from_tensor_slices((x_test, y_test))
3
4 train_dataset = train_dataset.shuffle(60000).batch(64)
5 test_dataset = test_dataset.batch(64)
```

In [6]:

```
1 loss_object = tf.keras.losses.SparseCategoricalCrossentropy()
2 optimizer = tf.keras.optimizers.Adam()
```

Create stateful metrics that can be used to accumulate values during training and logged at any point:

In [7]:

```
1 # Define our metrics
2 train_loss = tf.keras.metrics.Mean('train_loss', dtype=tf.float32)
3 train_accuracy = tf.keras.metrics.SparseCategoricalAccuracy('train_accuracy')
4 test_loss = tf.keras.metrics.Mean('test_loss', dtype=tf.float32)
5 test_accuracy = tf.keras.metrics.SparseCategoricalAccuracy('test_accuracy')
```

**Define the training and test functions:**

In [8]:



```
1 def train_step(model, optimizer, x_train, y_train):
2     with tf.GradientTape() as tape:
3         predictions = model(x_train, training=True)
4         loss = loss_object(y_train, predictions)
5         grads = tape.gradient(loss, model.trainable_variables)
6         optimizer.apply_gradients(zip(grads, model.trainable_variables))
7
8     train_loss(loss)
9     train_accuracy(y_train, predictions)
10
11 def test_step(model, x_test, y_test):
12     predictions = model(x_test)
13     loss = loss_object(y_test, predictions)
14
15     test_loss(loss)
16     test_accuracy(y_test, predictions)
```

Set up summary writers to write the summaries to disk in a different logs directory:

In [11]:



```
1 current_time = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
2 train_log_dir = 'logs/gradient_tape/' + current_time + '/train'
3 test_log_dir = 'logs/gradient_tape/' + current_time + '/test'
4 train_summary_writer = tf.summary.create_file_writer(train_log_dir)
5 test_summary_writer = tf.summary.create_file_writer(test_log_dir)
```

Start training. Use `tf.summary.scalar()` to log metrics (loss and accuracy) during training/testing within the scope of the summary writers to write the summaries to disk. You have control over which metrics to log and how often to do it. Other `tf.summary` functions enable logging other types of data.

In [13]:



```
1 EPOCHS = 5
2
3 for epoch in range(EPOCHS):
4     for (x_train, y_train) in train_dataset:
5         train_step(model, optimizer, x_train, y_train)
6     with train_summary_writer.as_default():
7         tf.summary.scalar('loss', train_loss.result(), step=epoch)
8         tf.summary.scalar('accuracy', train_accuracy.result(), step=epoch)
9
10    for (x_test, y_test) in test_dataset:
11        test_step(model, x_test, y_test)
12    with test_summary_writer.as_default():
13        tf.summary.scalar('loss', test_loss.result(), step=epoch)
14        tf.summary.scalar('accuracy', test_accuracy.result(), step=epoch)
15
16    template = 'Epoch {}, Loss: {}, Accuracy: {}, Test Loss: {}, Test Accur
17    print (template.format(epoch+1,
18                            train_loss.result(),
19                            train_accuracy.result()*100,
20                            test_loss.result(),
21                            test_accuracy.result()*100))
22
23    # Reset metrics every epoch
24    train_loss.reset_states()
25    test_loss.reset_states()
26    train_accuracy.reset_states()
27    test_accuracy.reset_states()
```

```
Epoch 1, Loss: 0.013252157717943192, Accuracy: 99.5533370971679
7, Test Loss: 0.06811141967773438, Test Accuracy: 98.32999420166
016
Epoch 2, Loss: 0.010231240652501583, Accuracy: 99.6533355712890
6, Test Loss: 0.06884602457284927, Test Accuracy: 98.22000122070
312
Epoch 3, Loss: 0.011641748249530792, Accuracy: 99.6116638183593
8, Test Loss: 0.06621949374675751, Test Accuracy: 98.40999603271
484
Epoch 4, Loss: 0.010976376011967659, Accuracy: 99.6333389282226
6, Test Loss: 0.07315755635499954, Test Accuracy: 98.29000091552
734
Epoch 5, Loss: 0.012019790709018707, Accuracy: 99.5849990844726
6, Test Loss: 0.07585285604000092, Test Accuracy: 98.25999450683
594
```

**To run TensorBoard, run the following command on Anaconda (Powershell) Prompt :**

```
tensorboard --logdir= path/to/log-directory
```

- For instance, `tensorboard --logdir logs/gradient_tape`

Connecting to `http://localhost:6006`

## 2. Tensorboard on Google Colab

- **TensorBoard on Colab** - TensorBoard.dev <https://tensorboard.dev/>  
(<https://tensorboard.dev/>)
  - [Example Code]:  
<https://colab.research.google.com/github/tensorflow/tensorboard/blob/master/docs/tbdev>  
(<https://colab.research.google.com/github/tensorflow/tensorboard/blob/master/docs/tbdev>)

