

# < Deep Learning - PART1 TF2 Basics >

## Ch 4. Workshop : FCDNs - MNIST for Digit Recognition

2021/10/01

[ Reference ] :

1. FRANÇOIS CHOLLET, **Deep Learning with Python**, Chapter 3, Section 5, Manning, 2018.  
(<https://tanhiamhuat.files.wordpress.com/2018/03/deeplearningwithpython.pdf>  
(<https://tanhiamhuat.files.wordpress.com/2018/03/deeplearningwithpython.pdf>))

2. TensorBoard - Guide : **Displaying image data in TensorBoard**,  
[https://www.tensorflow.org/tensorboard/image\\_summaries](https://www.tensorflow.org/tensorboard/image_summaries)  
([https://www.tensorflow.org/tensorboard/image\\_summaries](https://www.tensorflow.org/tensorboard/image_summaries)).

In [1]:



```
1 import tensorflow as tf
2 from datetime import datetime
3
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import io
7 import itertools
8 import sklearn
9 from sklearn.metrics import confusion_matrix
10
11 tf.__version__
```

Out[1]:

'2.4.1'

### Using Fully-Connected Deep Networks (FCDNs) to classify hand-written digits :

- In our case, we will configure our convnet to process inputs of size (28, 28, 1) , which is the format of **MNIST** images. We do this via passing the argument `input_shape=(28, 28, 1)` to our first hidden layer.
- The FCDN model consists of 2 hidden layers: the first hidden layer poccesses 128 hidden units, while the second one has 64 hidden units.
- There are 10 classes within the output layer.

In [2]:

```
1 from tensorflow.keras import models, layers
2
3 model = models.Sequential()
4 model.add(layers.Flatten(input_shape=(28, 28, 1)))
5 model.add(layers.Dense(128, activation='relu'))
6 model.add(layers.Dense(64, activation='relu'))
7 model.add(layers.Dense(10, activation='softmax'))
```

We are going to do 10-way classification, so we use a final layer with 10 outputs and a softmax activation. Now here's what our network looks like:

In [3]:

```
1 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 128)	100480
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 10)	650
Total params: 109,386		
Trainable params: 109,386		
Non-trainable params: 0		

As you can see, our (3, 3, 64) outputs were flattened into vectors of shape (576,) , before going through two Dense layers.

Now, let's train our convnet on the MNIST digits. We will reuse a lot of the code we have already covered in the MNIST example from Chapter 2.

In [4]:

```
1 mnist = tf.keras.datasets.mnist
2 (train_images, train_labels), (test_images, test_labels) = mnist.load_data(
3
4 # Reshape the datasets into 4D-Tensor format...
5 train_images = train_images.reshape((60000, 28, 28, 1))
6 train_images = train_images.astype('float32') / 255
7
8 test_images = test_images.reshape((10000, 28, 28, 1))
9 test_images = test_images.astype('float32') / 255
```

## [ About Loss Function ]: categorical\_crossentropy vs. sparse\_categorical\_crossentropy

- If the targets are **one-hot encoded**, use **categorical\_crossentropy** .

- Examples of one-hot encodings:

[1,0,0]

[0,1,0]

[0,0,1]

- But if the targets are **integers**, use **sparse\_categorical\_crossentropy** .

- Examples of integer encodings (for the sake of completion):

1

2

3

## Backpropagation

In [5]:



```
1 model.compile(optimizer='rmsprop',  
2               loss='sparse_categorical_crossentropy',  
3               metrics=['accuracy'])
```

## Output Data for TensorBoard

- Creating a File Directory for TensorBoard

In [6]:



```
1 # File-directory Path for TensorBoard  
2 log_dir="logs/fit/fcdn/" + datetime.now().strftime("%Y%m%d-%H%M%S")
```

- Output MNIST Images for TensorBoard

In [7]:



```
1 # tf.summary.create_file_writer() :
2 #   Creates a summary file writer for the given log directory.
3 log_images = log_dir + "/images"
4 file_writer = tf.summary.create_file_writer(log_images)
5
6 def plot_to_image(figure):
7     """Converts the matplotlib plot specified by 'figure'
8         to a PNG image and returns it. The supplied figure
9         is closed and inaccessible after this call."""
10    # Save the plot to a PNG in memory.
11    buf = io.BytesIO()
12    plt.savefig(buf, format='png')
13    # Closing the figure prevents it from being displayed
14    # directly inside the notebook.
15    plt.close(figure)
16    buf.seek(0)
17    # Convert PNG buffer to TF image
18    image = tf.image.decode_png(buf.getvalue(), channels=4)
19    # Add the batch dimension
20    image = tf.expand_dims(image, 0)
21    return image
22
23 def image_grid():
24     """Return a 5x5 grid of the MNIST images as a matplotlib figure."""
25     # Create a figure to contain the plot.
26     figure = plt.figure(figsize=(10,10))
27     for i in range(25):
28         # Start next subplot.
29         plt.subplot(5, 5, i + 1, title=class_names[train_labels[i]])
30         plt.xticks([])
31         plt.yticks([])
32         plt.grid(False)
33         plt.imshow(train_images[i], cmap=plt.cm.binary)
34     return figure
35
36 def log_image_plots(epoch, logs):
37     # Prepare the plot
38     figure = image_grid()
39     # Convert to image and Log
40     with file_writer.as_default():
41         tf.summary.image("Training data", plot_to_image(figure), step=0)
```

- Output Confusion Matrix for TensorBoard

In [8]:



```
1 def plot_confusion_matrix(cm, class_names):
2     """
3     Returns a matplotlib figure containing the plotted confusion matrix.
4
5     Args:
6         cm (array, shape = [n, n]): a confusion matrix of integer classes
7         class_names (array, shape = [n]): String names of the integer classes
8     """
9     figure = plt.figure(figsize=(8, 8))
10    plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
11    plt.title("Confusion matrix")
12    plt.colorbar()
13    tick_marks = np.arange(len(class_names))
14    plt.xticks(tick_marks, class_names, rotation=45)
15    plt.yticks(tick_marks, class_names)
16
17    # Compute the labels from the normalized confusion matrix.
18    labels = np.around(cm.astype('float') / cm.sum(axis=1)[:, np.newaxis],
19                       decimals=2)
20
21    # Use white text if squares are dark; otherwise black.
22    threshold = cm.max() / 2.
23    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
24        color = "white" if cm[i, j] > threshold else "black"
25        plt.text(j, i, labels[i, j], horizontalalignment="center", color=color)
26
27    plt.tight_layout()
28    plt.ylabel('True label')
29    plt.xlabel('Predicted label')
30    return figure
```

In [9]:



```
1 # tf.summary.create_file_writer() :
2 # Creates a summary file writer for the given log directory.
3 file_writer_cm = tf.summary.create_file_writer(log_dir + '/cm')
4
5 # Class names for MNIST digit images: 0 ~ 9
6 class_names = [str(i) for i in range(10)]
7
8 def log_confusion_matrix(epoch, logs):
9     # Use the model to predict the values from the validation dataset.
10    test_pred_raw = model.predict(test_images)
11    test_pred = np.argmax(test_pred_raw, axis=-1)
12
13    # Calculate the confusion matrix.
14    cm = sklearn.metrics.confusion_matrix(test_labels, test_pred)
15    # Log the confusion matrix as an image summary.
16    figure = plot_confusion_matrix(cm, class_names=class_names)
17    cm_image = plot_to_image(figure)
18
19    # Log the confusion matrix as an image summary.
20    with file_writer_cm.as_default():
21        tf.summary.image("Confusion Matrix", cm_image, step=epoch)
```

- **Callback methods**

In [10]:



```
1 # Define the Keras TensorBoard callback.
2 tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
3                                                         histogram_freq=1)
4
5 # Define the per-epoch callbacks for MNIST images.
6 image_callback = tf.keras.callbacks.LambdaCallback(on_epoch_end=log_image_p
7
8 # Define the per-epoch callbacks for confusion matrix.
9 cm_callback = tf.keras.callbacks.LambdaCallback(on_epoch_end=log_confusion_
```

In [11]:



```
1 # validation_split: Float between 0 and 1.
2 #           Fraction of the training data to be used as validation
3
4 history = model.fit(train_images,
5                     train_labels,
6                     epochs=15,
7                     batch_size=512,
8                     validation_split=0.1,
9                     callbacks=[tensorboard_callback,
10                              image_callback,
11                              cm_callback])
```

Epoch 1/15

106/106 [=====] - 2s 12ms/step - loss: 0.8763 - accuracy: 0.7545 - val\_loss: 0.2480 - val\_accuracy: 0.9265

Epoch 2/15

106/106 [=====] - 0s 4ms/step - loss: 0.2547 - accuracy: 0.9260 - val\_loss: 0.1697 - val\_accuracy: 0.9512

Epoch 3/15

106/106 [=====] - 0s 4ms/step - loss: 0.1849 - accuracy: 0.9457 - val\_loss: 0.1317 - val\_accuracy: 0.9630

Epoch 4/15

106/106 [=====] - 0s 4ms/step - loss: 0.1440 - accuracy: 0.9573 - val\_loss: 0.1181 - val\_accuracy: 0.9632

Epoch 5/15

106/106 [=====] - 0s 4ms/step - loss: 0.1132 - accuracy: 0.9666 - val\_loss: 0.0964 - val\_accuracy: 0.9710

Epoch 6/15

106/106 [=====] - 0s 4ms/step - loss: 0.0941 - accuracy: 0.9723 - val\_loss: 0.1071 - val\_accuracy: 0.9695

Epoch 7/15

106/106 [=====] - 0s 4ms/step - loss: 0.0784 - accuracy: 0.9765 - val\_loss: 0.0972 - val\_accuracy: 0.9715

Epoch 8/15

106/106 [=====] - 0s 4ms/step - loss: 0.0663 - accuracy: 0.9802 - val\_loss: 0.0825 - val\_accuracy: 0.9738

Epoch 9/15

106/106 [=====] - 0s 4ms/step - loss: 0.0576 - accuracy: 0.9831 - val\_loss: 0.0819 - val\_accuracy: 0.9767

Epoch 10/15

106/106 [=====] - 0s 4ms/step - loss: 0.0491 - accuracy: 0.9856 - val\_loss: 0.0863 - val\_accuracy: 0.9725

Epoch 11/15

106/106 [=====] - 0s 4ms/step - loss: 0.0408 - accuracy: 0.9880 - val\_loss: 0.0810 - val\_accuracy: 0.9775

Epoch 12/15

106/106 [=====] - 0s 4ms/step - loss: 0.0367 - accuracy: 0.9893 - val\_loss: 0.0856 - val\_accuracy: 0.9770

Epoch 13/15

106/106 [=====] - 0s 4ms/step - loss: 0.0322 - accuracy: 0.9915 - val\_loss: 0.0725 - val\_accuracy: 0.9790

Epoch 14/15

106/106 [=====] - 0s 4ms/step - loss: 0.0303 - accuracy: 0.9914 - val\_loss: 0.0788 - val\_accuracy: 0.9772

```
Epoch 15/15
106/106 [=====] - 0s 4ms/step - loss: 0.
0237 - accuracy: 0.9934 - val_loss: 0.0765 - val_accuracy: 0.9792
```

Let's evaluate the model on the test data:

In [12]:

```
1 # verbose: 0 or 1. Verbosity mode. 0 = silent, 1 = progress bar
2 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=0)
```

In [13]:

```
1 test_acc
```

Out[13]:

```
0.9771999716758728
```

In [14]:

```
1 test_loss
```

Out[14]:

```
0.07573146373033524
```

## Plotting results

- The call to `model.fit()` returns a `History` object.
- This object has a member `history`, which is a dictionary containing data about everything that happened during training.

Let's take a look at it:



In [15]:


```
1 history_dict = history.history
2 history_dict.keys()
```

Out[15]:

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

In [16]:


```
1 import matplotlib.pyplot as plt
2 %matplotlib inline
3
4 acc = history.history['accuracy']
5 val_acc = history.history['val_accuracy']
6 loss = history.history['loss']
7 val_loss = history.history['val_loss']
8
9 epochs = range(1, len(acc) + 1)
10
11 # "bo" is for "blue dot"
12 plt.plot(epochs, loss, 'bo', label='Training loss')
13 # b is for "solid blue line"
14 plt.plot(epochs, val_loss, 'b', label='Validation loss')
15 plt.title('Training and validation loss')
16 plt.xlabel('Epochs')
17 plt.ylabel('Loss')
18 plt.legend()
19
20 plt.show()
```



In [17]:



```
1 plt.clf() # clear figure
2 acc_values = history_dict['accuracy']
3 val_acc_values = history_dict['val_accuracy']
4
5 plt.plot(epochs, acc, 'bo', label='Training acc')
6 plt.plot(epochs, val_acc, 'b', label='Validation acc')
7 plt.title('Training and validation accuracy')
8 plt.xlabel('Epochs')
9 plt.ylabel('Accuracy')
10 plt.legend()
11
12 plt.show()
```



## Prediction

In [18]:



```
1 test_predict = model.predict(test_images)
2 test_predict
```

Out[18]:

```
array([[1.5047672e-07, 1.2711048e-09, 3.3050892e-05, ..., 9.99353
59e-01,
        1.1711770e-06, 2.3033490e-05],
       [4.6101580e-09, 1.4413783e-05, 9.9997699e-01, ..., 5.35253
77e-14,
        2.5796912e-06, 7.9698786e-13],
       [8.3861751e-06, 9.9881315e-01, 4.3890518e-04, ..., 3.54902
09e-04,
        1.8984842e-04, 1.4843447e-06],
       ...,
       [1.4883571e-13, 1.1322564e-13, 1.0123786e-13, ..., 3.14632
59e-08,
        3.3519623e-08, 7.8649146e-07],
       [6.2846061e-12, 2.0362026e-13, 2.0849162e-12, ..., 8.51185
44e-11,
        2.1644304e-05, 1.7771976e-12],
       [3.4703443e-10, 3.5609900e-15, 9.8466929e-12, ..., 1.90469
85e-18,
        2.1573755e-11, 6.8828145e-11]], dtype=float32)
```

In [19]:



```
1 import numpy as np
2 test_predict_result = np.array([np.argmax(test_predict[i]) for i in range(1
3 test_predict_result
```

Out[19]:

```
array([7, 2, 1, ..., 4, 5, 6], dtype=int64)
```

In [20]:



```
1 test_labels
```

Out[20]:


```
array([7, 2, 1, ..., 4, 5, 6], dtype=uint8)
```

## Confusion Matrix

In [21]:



```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 sns.set()
4 %matplotlib inline
5
6 from sklearn.metrics import confusion_matrix
7
8 mat = confusion_matrix(test_labels, test_predict_result)
9
10 plt.figure(figsize=(10,8))
11 sns.heatmap(mat, square=False, annot=True, fmt='d', cbar=True)
12 plt.xlim((0, 10))
13 plt.ylim((10, 0))
14 plt.xlabel('predicted value')
15 plt.ylabel('true value')
16 plt.show()
17
18 mat
```



Out[21]:

```
array([[ 970,    0,    0,    0,    2,    0,    1,    1,    4,
```

```

2],
0],
0],
5],
4],
2],
0],
9],
4],
3]],
dtype=int64)

```

## TensorBoard

---

**To run TensorBoard, run the following command on Anaconda Prompt :**

```
tensorboard --logdir= path/to/log-directory
```

- For instance, `tensorboard --logdir logs/fit/fcdn`

Connecting to `http://localhost:6006`

---