

< Deep Learning - PART2 TF2 CNNs >

Ch 5. CNNs Workshop 3 - CIFAR10 : Image Classifier with Data Augmentation & Dropout

2021/10/01

[Reference] :

- Keras Documentation : **Train a simple deep CNN on the CIFAR10 small images dataset**
https://keras.io/examples/cifar10_cnn/ (https://keras.io/examples/cifar10_cnn/).
- **CIFAR-10 and CIFAR-100 datasets** <https://www.cs.toronto.edu/~kriz/cifar.html>
(<https://www.cs.toronto.edu/~kriz/cifar.html>)
- 李飛飛教授 : Convolutional Neural Networks (CNNs / ConvNets)
(<https://cs231n.github.io/convolutional-networks/> (<https://cs231n.github.io/convolutional-networks/>))
- `tf.keras.layers.Conv2D`
(https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D
(https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D))

[NOTE] : Running the following code on Google Colab with GPU !!

Running TensorFlow 2.0 on Google Colab

In [2]:



```
1 from __future__ import print_function
2
3 import tensorflow as tf
4 print(tf.__version__)
5
6 from tensorflow.keras.datasets import cifar10
7 from tensorflow.keras.preprocessing.image import ImageDataGenerator
8 from tensorflow.keras.models import Sequential
9 from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten
10 from tensorflow.keras.layers import Conv2D, MaxPooling2D
11 import os
```

2.1.0

Hyperparameters

In [0]:



```
1 batch_size = 32
2 num_classes = 10
3 epochs = 50
4 data_augmentation = True
5 num_predictions = 20
6
7 save_dir = os.path.join(os.getcwd(), 'saved_models')
8 model_name = 'keras_cifar10_trained_model.h5'
```

Loading data...

In [4]:



```
1 # The data, split between train and test sets:
2 (x_train, y_train), (x_test, y_test) = cifar10.load_data()
3 print('x_train shape:', x_train.shape)
4 print(x_train.shape[0], 'train samples')
5 print(x_test.shape[0], 'test samples')
6
7 # Convert class vectors to binary class matrices.
8 y_train = tf.keras.utils.to_categorical(y_train, num_classes)
9 y_test = tf.keras.utils.to_categorical(y_test, num_classes)
```

```
x_train shape: (50000, 32, 32, 3)
50000 train samples
10000 test samples
```

Forward Propagation

- Building a training model with `tf.keras`

In [0]:



```
1 model = Sequential()
2 model.add(Conv2D(32, (3, 3), padding='same',
3                 input_shape=x_train.shape[1:]))
4 model.add(Activation('relu'))
5 model.add(Conv2D(32, (3, 3)))
6 model.add(Activation('relu'))
7 model.add(MaxPooling2D(pool_size=(2, 2)))
8 model.add(Dropout(0.25))
9
10 model.add(Conv2D(64, (3, 3), padding='same'))
11 model.add(Activation('relu'))
12 model.add(Conv2D(64, (3, 3)))
13 model.add(Activation('relu'))
14 model.add(MaxPooling2D(pool_size=(2, 2)))
15 model.add(Dropout(0.25))
16
17 model.add(Flatten())
18 model.add(Dense(512))
19 model.add(Activation('relu'))
20 model.add(Dropout(0.5))
21 model.add(Dense(num_classes))
22 model.add(Activation('softmax'))
```

Backpropagation

In [0]:



```
1 # initiate RMSprop optimizer
2 opt = tf.keras.optimizers.RMSprop(learning_rate=0.0001, decay=1e-6)
3
4 # Let's train the model using RMSprop
5 model.compile(loss='categorical_crossentropy',
6               optimizer=opt,
7               metrics=['accuracy'])
```

Normalization

In [0]:



```
1 x_train = x_train.astype('float32')
2 x_test = x_test.astype('float32')
3 x_train /= 255
4 x_test /= 255
```

Data Augmentation

In [8]:



```
1 if not data_augmentation:
2     print('Not using data augmentation.')
3     model.fit(x_train, y_train,
4               batch_size=batch_size,
5               epochs=epochs,
6               validation_data=(x_test, y_test),
7               shuffle=True)
8 else:
9     print('Using real-time data augmentation.')
10    # This will do preprocessing and realtime data augmentation:
11    datagen = ImageDataGenerator(
12        featurewise_center=False, # set input mean to 0 over the dataset
13        samplewise_center=False, # set each sample mean to 0
14        featurewise_std_normalization=False, # divide inputs by std of the
15        samplewise_std_normalization=False, # divide each input by its std
16        zca_whitening=False, # apply ZCA whitening
17        zca_epsilon=1e-06, # epsilon for ZCA whitening
18        rotation_range=0, # randomly rotate images in the range (degrees,
19        # randomly shift images horizontally (fraction of total width)
20        width_shift_range=0.1,
21        # randomly shift images vertically (fraction of total height)
22        height_shift_range=0.1,
23        shear_range=0., # set range for random shear
24        zoom_range=0., # set range for random zoom
25        channel_shift_range=0., # set range for random channel shifts
26        # set mode for filling points outside the input boundaries
27        fill_mode='nearest',
28        cval=0., # value used for fill_mode = "constant"
29        horizontal_flip=True, # randomly flip images
30        vertical_flip=False, # randomly flip images
31        # set rescaling factor (applied before any other transformation)
32        rescale=None,
33        # set function that will be applied on each input
34        preprocessing_function=None,
35        # image data format, either "channels_first" or "channels_last"
36        data_format=None,
37        # fraction of images reserved for validation (strictly between 0 and
38        validation_split=0.0)
39
40    # Compute quantities required for feature-wise normalization
41    # (std, mean, and principal components if ZCA whitening is applied).
42    datagen.fit(x_train)
43
44    # Fit the model on the batches generated by datagen.flow().
45    model.fit_generator(datagen.flow(x_train, y_train,
46                                    batch_size=batch_size),
47                        epochs=epochs,
48                        validation_data=(x_test, y_test),
49                        workers=4)
```

Using real-time data augmentation.

WARNING:tensorflow:From <ipython-input-8-58db1f5097ea>:49: Model.
fit_generator (from tensorflow.python.keras.engine.training) is d
eprecated and will be removed in a future version.

Instructions for updating:

Please use `Model.fit`, which supports generators.

WARNING:tensorflow:sample_weight modes were coerced from

...

to

['...']

Train for 1563 steps, validate on 10000 samples

Epoch 1/50

1563/1563 [=====] - 34s 22ms/step - loss: 1.8522 - accuracy: 0.3156 - val_loss: 1.5686 - val_accuracy: 0.4282

Epoch 2/50

1563/1563 [=====] - 31s 20ms/step - loss: 1.5820 - accuracy: 0.4212 - val_loss: 1.3962 - val_accuracy: 0.4895

Epoch 3/50

1563/1563 [=====] - 31s 20ms/step - loss: 1.4621 - accuracy: 0.4668 - val_loss: 1.2855 - val_accuracy: 0.5424

Epoch 4/50

1563/1563 [=====] - 31s 20ms/step - loss: 1.3848 - accuracy: 0.5029 - val_loss: 1.3087 - val_accuracy: 0.5283

Epoch 5/50

1563/1563 [=====] - 31s 20ms/step - loss: 1.3199 - accuracy: 0.5261 - val_loss: 1.1952 - val_accuracy: 0.5807

Epoch 6/50

1563/1563 [=====] - 31s 20ms/step - loss: 1.2646 - accuracy: 0.5492 - val_loss: 1.0823 - val_accuracy: 0.6243

Epoch 7/50

1563/1563 [=====] - 32s 20ms/step - loss: 1.2151 - accuracy: 0.5705 - val_loss: 1.0979 - val_accuracy: 0.6160

Epoch 8/50

1563/1563 [=====] - 31s 20ms/step - loss: 1.1746 - accuracy: 0.5853 - val_loss: 1.0508 - val_accuracy: 0.6367

Epoch 9/50

1563/1563 [=====] - 31s 20ms/step - loss: 1.1351 - accuracy: 0.6001 - val_loss: 0.9939 - val_accuracy: 0.6555

Epoch 10/50

1563/1563 [=====] - 31s 20ms/step - loss: 1.1063 - accuracy: 0.6082 - val_loss: 1.0137 - val_accuracy: 0.6450

Epoch 11/50

1563/1563 [=====] - 32s 20ms/step - loss: 1.0778 - accuracy: 0.6211 - val_loss: 0.9684 - val_accuracy: 0.6643

Epoch 12/50

1563/1563 [=====] - 31s 20ms/step - loss: 1.0528 - accuracy: 0.6297 - val_loss: 0.9301 - val_accuracy: 0.6730

Epoch 13/50

1563/1563 [=====] - 32s 20ms/step - loss

s: 1.0243 - accuracy: 0.6415 - val_loss: 0.9945 - val_accuracy: 0.6649
Epoch 14/50
1563/1563 [=====] - 31s 20ms/step - loss: 1.0060 - accuracy: 0.6457 - val_loss: 0.8892 - val_accuracy: 0.6916
Epoch 15/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.9801 - accuracy: 0.6570 - val_loss: 0.8424 - val_accuracy: 0.7071
Epoch 16/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.9681 - accuracy: 0.6589 - val_loss: 0.8317 - val_accuracy: 0.7079
Epoch 17/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.9542 - accuracy: 0.6675 - val_loss: 0.8605 - val_accuracy: 0.6934
Epoch 18/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.9347 - accuracy: 0.6730 - val_loss: 0.8303 - val_accuracy: 0.7125
Epoch 19/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.9208 - accuracy: 0.6768 - val_loss: 0.7733 - val_accuracy: 0.7314
Epoch 20/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.9115 - accuracy: 0.6806 - val_loss: 0.8550 - val_accuracy: 0.7100
Epoch 21/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.9024 - accuracy: 0.6855 - val_loss: 0.7837 - val_accuracy: 0.7230
Epoch 22/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8937 - accuracy: 0.6900 - val_loss: 0.7865 - val_accuracy: 0.7286
Epoch 23/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8833 - accuracy: 0.6945 - val_loss: 0.7812 - val_accuracy: 0.7257
Epoch 24/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8740 - accuracy: 0.6971 - val_loss: 0.7955 - val_accuracy: 0.7253
Epoch 25/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8716 - accuracy: 0.6961 - val_loss: 0.7891 - val_accuracy: 0.7319
Epoch 26/50
1563/1563 [=====] - 32s 20ms/step - loss: 0.8597 - accuracy: 0.7020 - val_loss: 0.7382 - val_accuracy: 0.7421
Epoch 27/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8568 - accuracy: 0.7059 - val_loss: 0.7289 - val_accuracy:

0.7510
Epoch 28/50
1563/1563 [=====] - 32s 20ms/step - loss: 0.8541 - accuracy: 0.7052 - val_loss: 0.7725 - val_accuracy: 0.7366
Epoch 29/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8476 - accuracy: 0.7087 - val_loss: 0.7659 - val_accuracy: 0.7441
Epoch 30/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8410 - accuracy: 0.7123 - val_loss: 0.7346 - val_accuracy: 0.7480
Epoch 31/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8377 - accuracy: 0.7139 - val_loss: 0.7613 - val_accuracy: 0.7409
Epoch 32/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8360 - accuracy: 0.7144 - val_loss: 0.7264 - val_accuracy: 0.7492
Epoch 33/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8283 - accuracy: 0.7155 - val_loss: 0.7821 - val_accuracy: 0.7365
Epoch 34/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8243 - accuracy: 0.7176 - val_loss: 0.7258 - val_accuracy: 0.7615
Epoch 35/50
1563/1563 [=====] - 31s 20ms/step - loss: 0.8299 - accuracy: 0.7177 - val_loss: 0.7439 - val_accuracy: 0.7461
Epoch 36/50
1563/1563 [=====] - 33s 21ms/step - loss: 0.8227 - accuracy: 0.7201 - val_loss: 0.7150 - val_accuracy: 0.7599
Epoch 37/50
1563/1563 [=====] - 32s 20ms/step - loss: 0.8220 - accuracy: 0.7201 - val_loss: 0.7034 - val_accuracy: 0.7635
Epoch 38/50
1563/1563 [=====] - 32s 20ms/step - loss: 0.8213 - accuracy: 0.7209 - val_loss: 0.7193 - val_accuracy: 0.7582
Epoch 39/50
1563/1563 [=====] - 32s 20ms/step - loss: 0.8152 - accuracy: 0.7207 - val_loss: 0.7304 - val_accuracy: 0.7671
Epoch 40/50
1563/1563 [=====] - 32s 20ms/step - loss: 0.8172 - accuracy: 0.7227 - val_loss: 0.7036 - val_accuracy: 0.7684
Epoch 41/50
1563/1563 [=====] - 32s 20ms/step - loss: 0.8120 - accuracy: 0.7248 - val_loss: 0.6881 - val_accuracy: 0.7684

```

Epoch 42/50
1563/1563 [=====] - 31s 20ms/step - los
s: 0.8116 - accuracy: 0.7257 - val_loss: 0.7248 - val_accuracy:
0.7570
Epoch 43/50
1563/1563 [=====] - 32s 20ms/step - los
s: 0.8085 - accuracy: 0.7258 - val_loss: 0.7283 - val_accuracy:
0.7587
Epoch 44/50
1563/1563 [=====] - 31s 20ms/step - los
s: 0.8088 - accuracy: 0.7257 - val_loss: 0.7152 - val_accuracy:
0.7633
Epoch 45/50
1563/1563 [=====] - 32s 20ms/step - los
s: 0.8068 - accuracy: 0.7241 - val_loss: 0.7186 - val_accuracy:
0.7553
Epoch 46/50
1563/1563 [=====] - 32s 20ms/step - los
s: 0.8049 - accuracy: 0.7264 - val_loss: 0.7038 - val_accuracy:
0.7662
Epoch 47/50
1563/1563 [=====] - 32s 20ms/step - los
s: 0.8063 - accuracy: 0.7248 - val_loss: 0.7500 - val_accuracy:
0.7517
Epoch 48/50
1563/1563 [=====] - 32s 20ms/step - los
s: 0.7995 - accuracy: 0.7315 - val_loss: 0.6898 - val_accuracy:
0.7786
Epoch 49/50
1563/1563 [=====] - 33s 21ms/step - los
s: 0.8005 - accuracy: 0.7287 - val_loss: 0.7452 - val_accuracy:
0.7496
Epoch 50/50
1563/1563 [=====] - 32s 21ms/step - los
s: 0.8020 - accuracy: 0.7296 - val_loss: 0.7083 - val_accuracy:
0.7645

```

Saving the model...

In [9]:



```

1  # Save model and weights
2  if not os.path.isdir(save_dir):
3      os.makedirs(save_dir)
4  model_path = os.path.join(save_dir, model_name)
5  model.save(model_path)
6  print('Saved trained model at %s ' % model_path)

```

Saved trained model at /content/saved_models/keras_cifar10_traine
d_model.h5

Evaluation

In [10]:



```
1 # Score trained model.
2 scores = model.evaluate(x_test, y_test, verbose=2)
3 print('Test loss:', scores[0])
4 print('Test accuracy:', scores[1])
```

10000/10000 - 1s - loss: 0.7083 - accuracy: 0.7645

Test loss: 0.708253707408905

Test accuracy: 0.7645