

# < Deep Learning - PART2 TF2 CNNs >

## Ch 5. CNNs Workshop 4 - CIFAR10 : Image Classifier on Google Colab & Google Drive ¶

2021/10/01

### [ Reference ] :

- TensorFlow Core - Tutorials: **Convolutional Neural Network (CNN)**  
[https://www.tensorflow.org/tutorials/images/cnn?hl=zh\\_tw](https://www.tensorflow.org/tutorials/images/cnn?hl=zh_tw)  
([https://www.tensorflow.org/tutorials/images/cnn?hl=zh\\_tw](https://www.tensorflow.org/tutorials/images/cnn?hl=zh_tw))
- **CIFAR-10 and CIFAR-100 datasets** <https://www.cs.toronto.edu/~kriz/cifar.html>  
(<https://www.cs.toronto.edu/~kriz/cifar.html>)
- 李飛飛教授 : Convolutional Neural Networks (教學投影片)  
([http://cs231n.stanford.edu/slides/2017/cs231n\\_2017\\_lecture5.pdf](http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture5.pdf)  
([http://cs231n.stanford.edu/slides/2017/cs231n\\_2017\\_lecture5.pdf](http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture5.pdf)))
- 李飛飛教授 : Convolutional Neural Networks (CNNs / ConvNets)  
(<https://cs231n.github.io/convolutional-networks/>) (<https://cs231n.github.io/convolutional-networks/>)
- `tf.keras.layers.Conv2D`  
([https://www.tensorflow.org/api\\_docs/python/tf/keras/layers/Conv2D](https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D)  
([https://www.tensorflow.org/api\\_docs/python/tf/keras/layers/Conv2D](https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D)))

In [ ]:

```
1 import tensorflow as tf
2 print(tf.__version__)
3
4 from tensorflow.keras import layers, models
```

2.3.0

### Download CIFAR-10 Python Version dataset from

- <https://www.cs.toronto.edu/~kriz/cifar.html>  
(<https://www.cs.toronto.edu/~kriz/cifar.html>)
- The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.

### [ Reference ] :

- Learning Multiple Layers of Features from Tiny Images , Alex Krizhevsky, 2009.  
<https://www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf>

(<https://www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf>)

In [ ]:

```
1 import numpy as np
2
3 class CifarLoader(object):
4     def __init__(self, source_files):
5         self._source = source_files
6         self._i = 0
7         self.images = None
8         self.labels = None
9
10    def load(self):
11        data = [unpickle(f) for f in self._source]
12        # print(data)
13        images = np.vstack([d[b"data"] for d in data])
14        n = len(images) # 32 x 32 x 3 channels
15        self.images = images.reshape(n, 3, 32, 32).transpose(0, 2, 3, 1).as
16        self.labels = one_hot(np.hstack([d[b"labels"] for d in data]), 10)
17        return self
18
19    def next_batch(self, batch_size):
20        x, y = self.images[self._i : self._i+batch_size], self.labels[self.
21        self._i = (self._i + batch_size) % len(self.images)
22        return x, y
```

## Loading dataset and running this program on the Google's Colab...

In [ ]:

```
1 # Load the Drive helper and mount
2 from google.colab import drive
3
4 # This will prompt for authorization.
5 drive.mount('/content/drive')
6
7 ## Go to this URL in a browser: https://accounts.google.com/o/oauth2/.....
8 ## Enter your authorization code:
9 ## .....
10 ## Mounted at /content/drive
```

Mounted at /content/drive

In [ ]:



```
1 # After executing the cell above, Drive
2 # files will be present in "/content/drive/My Drive".
3 # !ls "/content/drive/My Drive"
4 !ls "/content/drive/My Drive/Colab Notebooks/cifar-10-batches-py"
5 ## Loading CIFAR-10 dataset from your Google Drive
6 DATA_PATH = "/content/drive/My Drive/Colab Notebooks/cifar-10-batches-py"
```

```
batches.meta  data_batch_2  data_batch_4  readme.html
data_batch_1  data_batch_3  data_batch_5  test_batch
```

## Loading dataset from the unzipped files...

In [ ]:



```
1 import pickle
2 import os
3
4 ## Loading CIFAR-10 dataset from your Google Drive
5 DATA_PATH = "/content/drive/My Drive/Colab Notebooks/cifar-10-batches-py"
6
7 ## < for CIFAR-10 dataset on local host >
8 ## DATA_PATH = "./cifar-10-batches-py"
9
10 def unpickle(file):
11     with open(os.path.join(DATA_PATH, file), 'rb') as fo:
12         dict = pickle.load(fo, encoding='bytes') ## encoding='bytes'
13     return dict
14
15 def one_hot(vec, vals=10):
16     n = len(vec)
17     out = np.zeros((n, vals))
18     out[range(n), vec] = 1
19     return out
```

In [ ]:




```
1 class CifarDataManager(object):
2     def __init__(self):
3         self.train = CifarLoader(["data_batch_{}".format(i)
4                                   for i in range(1, 6)]).load()
5         self.test = CifarLoader(["test_batch"]).load()
```

In [ ]:



```
1 import matplotlib.pyplot as plt
2 %matplotlib inline
3
4 def display_cifar(images, size):
5     n = len(images)
6     plt.figure()
7     plt.gca().set_axis_off()
8     im = np.vstack([np.hstack([images[np.random.choice(n)] for i in range(s
9         for i in range(size)])
10    plt.imshow(im)
11    plt.show()
12
13 cifar = CifarDataManager() # Loading CIFAR-10 dataset
14 images = cifar.train.images
15 display_cifar(images, 10)
```



In [ ]:



```
1 print("Number of train images: {}".format(len(cifar.train.images)))
2 print("Number of train labels: {}".format(len(cifar.train.labels)))
3 print("Number of test images: {}".format(len(cifar.test.images)))
4 print("Number of test images: {}".format(len(cifar.test.labels)))
```

```
Number of train images: 50000
Number of train labels: 50000
Number of test images: 10000
Number of test images: 10000
```

In [ ]:



```
1 print("Train images: {}".format(cifar.train.images.shape))
2 print("Train labels: {}".format(cifar.train.labels.shape))
3 print(" Test images: {}".format(cifar.test.images.shape))
4 print(" Test labels: {}".format(cifar.test.labels.shape))
```

```
Train images: (50000, 32, 32, 3)
Train labels: (50000, 10)
Test images: (10000, 32, 32, 3)
Test labels: (10000, 10)
```

---

## < CNN Model > :

---

- **Forward Propagation**

In [200]:



```
1 model = models.Sequential()
2
3 # 1st (Conv Layer + Batch Norm) * 3 + MaxPooling Layer
4 model.add(layers.Conv2D(32, (3, 3), activation='relu',
5 padding='same', input_shape=(32, 32, 3)))
6 model.add(layers.BatchNormalization()) # Batch Norm
7 model.add(layers.Conv2D(32, (3, 3), activation='relu', padding='same'))
8 model.add(layers.BatchNormalization()) # Batch Norm
9 model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
10 model.add(layers.BatchNormalization()) # Batch Norm
11 model.add(layers.MaxPooling2D((2, 2)))
12 model.add(layers.Dropout(0.25))
13
14 # 2nd (Conv Layer + Batch Norm) * 2 + MaxPooling Layer
15 model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
16 model.add(layers.BatchNormalization()) # Batch Norm
17 model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
18 model.add(layers.BatchNormalization()) # Batch Norm
19 model.add(layers.MaxPooling2D((2, 2)))
20 model.add(layers.Dropout(0.25))
21
22 # 3rd (Conv Layer + Batch Norm) * 2 + MaxPooling Layer
23 model.add(layers.Conv2D(128, (3, 3), activation='relu', padding='same'))
24 model.add(layers.BatchNormalization()) # Batch Norm
25 model.add(layers.Conv2D(128, (3, 3), activation='relu', padding='same'))
26 model.add(layers.BatchNormalization()) # Batch Norm
27 model.add(layers.MaxPooling2D((2, 2)))
28 model.add(layers.Dropout(0.5))
29
30 # 4th (Conv Layer + Batch Norm) * 2
31 model.add(layers.Conv2D(256, (3, 3), activation='relu', padding='same'))
32 model.add(layers.BatchNormalization()) # Batch Norm
33 model.add(layers.Conv2D(256, (3, 3), activation='relu', padding='same'))
34 model.add(layers.BatchNormalization()) # Batch Norm
35 model.add(layers.MaxPooling2D((2, 2)))
36 model.add(layers.Dropout(0.5))
```

In [201]:



```
1 model.add(layers.Flatten())
2 model.add(layers.Dense(128, activation='relu'))
3 model.add(layers.BatchNormalization()) # Batch Norm
4 model.add(layers.Dropout(0.25))
5 model.add(layers.Dense(10, activation='softmax'))
```

In [202]:



```
1 model.summary()
```

Model: "sequential\_30"

Layer (type)	Output Shape	Param #
conv2d_245 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_280 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_246 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_281 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_247 (Conv2D)	(None, 32, 32, 64)	18496
batch_normalization_282 (Batch Normalization)	(None, 32, 32, 64)	256
max_pooling2d_111 (MaxPooling2D)	(None, 16, 16, 64)	0
dropout_149 (Dropout)	(None, 16, 16, 64)	0
conv2d_248 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_283 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_249 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_284 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_112 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_150 (Dropout)	(None, 8, 8, 64)	0
conv2d_250 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_285 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_251 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_286 (Batch Normalization)	(None, 8, 8, 128)	512
max_pooling2d_113 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_151 (Dropout)	(None, 4, 4, 128)	0
conv2d_252 (Conv2D)	(None, 4, 4, 256)	295168
batch_normalization_287 (Batch Normalization)	(None, 4, 4, 256)	1024
conv2d_253 (Conv2D)	(None, 4, 4, 256)	590080
batch_normalization_288 (Batch Normalization)	(None, 4, 4, 256)	1024

max_pooling2d_114 (MaxPoolin	(None, 2, 2, 256)	0
dropout_152 (Dropout)	(None, 2, 2, 256)	0
flatten_29 (Flatten)	(None, 1024)	0
dense_64 (Dense)	(None, 128)	131200
batch_normalization_289 (Bat	(None, 128)	512
dropout_153 (Dropout)	(None, 128)	0
dense_65 (Dense)	(None, 10)	1290
=====		
Total params: 1,346,282		
Trainable params: 1,343,978		
Non-trainable params: 2,304		

## • Back-Propagation

In [203]:



```
1 model.compile(optimizer='adam',
2               loss='categorical_crossentropy',
3               metrics=['accuracy'])
```

## Training ...

- BATCH\_SIZE = 32 (default value : 32)



In [204]:



```
1 history = model.fit(cifar.train.images, cifar.train.labels, epochs=30,  
2                     validation_split=0.1)
```

Epoch 1/30

1407/1407 [=====] - 36s 26ms/step - loss: 1.6297 - accuracy: 0.4180 - val\_loss: 1.1772 - val\_accuracy: 0.5842

Epoch 2/30

1407/1407 [=====] - 36s 26ms/step - loss: 1.0685 - accuracy: 0.6246 - val\_loss: 0.9373 - val\_accuracy: 0.6838

Epoch 3/30

1407/1407 [=====] - 36s 25ms/step - loss: 0.8717 - accuracy: 0.6985 - val\_loss: 0.7267 - val\_accuracy: 0.7464

Epoch 4/30

1407/1407 [=====] - 36s 25ms/step - loss: 0.7530 - accuracy: 0.7415 - val\_loss: 0.7063 - val\_accuracy: 0.7534

Epoch 5/30

1407/1407 [=====] - 36s 25ms/step - loss: 0.6714 - accuracy: 0.7692 - val\_loss: 0.6229 - val\_accuracy: 0.7980

Epoch 6/30

1407/1407 [=====] - 36s 25ms/step - loss: 0.6237 - accuracy: 0.7879 - val\_loss: 0.6122 - val\_accuracy: 0.7910

Epoch 7/30

1407/1407 [=====] - 36s 25ms/step - loss: 0.5640 - accuracy: 0.8100 - val\_loss: 0.5447 - val\_accuracy: 0.8166

Epoch 8/30

1407/1407 [=====] - 36s 25ms/step - loss: 0.5254 - accuracy: 0.8217 - val\_loss: 0.6024 - val\_accuracy: 0.7990

Epoch 9/30

1407/1407 [=====] - 36s 25ms/step - loss: 0.4881 - accuracy: 0.8335 - val\_loss: 0.5462 - val\_accuracy: 0.8170

Epoch 10/30

1407/1407 [=====] - 36s 25ms/step - loss: 0.4608 - accuracy: 0.8414 - val\_loss: 0.5271 - val\_accuracy: 0.8266

Epoch 11/30

1407/1407 [=====] - 36s 26ms/step - loss: 0.4340 - accuracy: 0.8521 - val\_loss: 0.4729 - val\_accuracy: 0.8408

Epoch 12/30

1407/1407 [=====] - 36s 25ms/step - loss: 0.4074 - accuracy: 0.8621 - val\_loss: 0.4527 - val\_accuracy: 0.8502

Epoch 13/30

1407/1407 [=====] - 36s 25ms/step - loss: 0.3886 - accuracy: 0.8671 - val\_loss: 0.4408 - val\_accuracy:

0.8558  
Epoch 14/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.3676 - accuracy: 0.8753 - val\_loss: 0.4486 - val\_accuracy: 0.8522  
Epoch 15/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.3516 - accuracy: 0.8808 - val\_loss: 0.4873 - val\_accuracy: 0.8440  
Epoch 16/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.3299 - accuracy: 0.8876 - val\_loss: 0.4335 - val\_accuracy: 0.8630  
Epoch 17/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.3204 - accuracy: 0.8894 - val\_loss: 0.4237 - val\_accuracy: 0.8660  
Epoch 18/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.3115 - accuracy: 0.8938 - val\_loss: 0.4284 - val\_accuracy: 0.8658  
Epoch 19/30  
1407/1407 [=====] - 36s 26ms/step - loss: 0.3003 - accuracy: 0.8970 - val\_loss: 0.4308 - val\_accuracy: 0.8640  
Epoch 20/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.2895 - accuracy: 0.8990 - val\_loss: 0.4123 - val\_accuracy: 0.8666  
Epoch 21/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.2723 - accuracy: 0.9065 - val\_loss: 0.4504 - val\_accuracy: 0.8614  
Epoch 22/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.2673 - accuracy: 0.9094 - val\_loss: 0.4262 - val\_accuracy: 0.8622  
Epoch 23/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.2600 - accuracy: 0.9108 - val\_loss: 0.4090 - val\_accuracy: 0.8734  
Epoch 24/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.2544 - accuracy: 0.9116 - val\_loss: 0.4136 - val\_accuracy: 0.8788  
Epoch 25/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.2436 - accuracy: 0.9158 - val\_loss: 0.4543 - val\_accuracy: 0.8686  
Epoch 26/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.2408 - accuracy: 0.9156 - val\_loss: 0.3975 - val\_accuracy: 0.8820  
Epoch 27/30  
1407/1407 [=====] - 36s 25ms/step - loss: 0.2309 - accuracy: 0.9218 - val\_loss: 0.4165 - val\_accuracy: 0.8758

```
Epoch 28/30
1407/1407 [=====] - 36s 26ms/step - loss: 0.2229 - accuracy: 0.9233 - val_loss: 0.4247 - val_accuracy: 0.8750
Epoch 29/30
1407/1407 [=====] - 36s 25ms/step - loss: 0.2180 - accuracy: 0.9253 - val_loss: 0.4129 - val_accuracy: 0.8808
Epoch 30/30
1407/1407 [=====] - 36s 25ms/step - loss: 0.2110 - accuracy: 0.9272 - val_loss: 0.4270 - val_accuracy: 0.8796
```

## Evaluation

In [205]:



```
1 test_loss, test_acc = model.evaluate(x=cifar.test.images, y=cifar.test.labels)
2 test_loss, test_acc
```

```
313/313 [=====] - 3s 9ms/step - loss: 0.4396 - accuracy: 0.8725
```

Out[205]:

```
(0.4396136701107025, 0.8725000023841858)
```

In [206]:




```
1 history.history.keys()
```

Out[206]:

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```


In [211]:

```
1 plt.plot(history.history['loss'], label='loss')
2 plt.plot(history.history['val_loss'], label = 'val_loss')
3 plt.xlabel('Epoch')
4 plt.ylabel('Loss')
5 plt.ylim([0, 1.5])
6 plt.legend(loc='lower left')
7 plt.show()
```



In [212]:

```
1 plt.plot(history.history['accuracy'], label='accuracy')
2 plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
3 plt.xlabel('Epoch')
4 plt.ylabel('Accuracy')
5 plt.ylim([0.3, 1])
6 plt.legend(loc='lower right')
7 plt.show()
```



## Confusion Matrix

In [213]:



```
1 test_predict = model.predict(cifar.test.images)
2 test_predict
```

Out[213]:

```
array([[1.4182602e-07, 6.2875465e-08, 1.6014928e-06, ..., 1.25194
69e-07,
        8.0835086e-08, 4.2166075e-08],
       [1.9468818e-05, 1.6829812e-03, 4.2885876e-08, ..., 8.36456
15e-09,
        9.9829668e-01, 1.8060865e-07],
       [2.0425217e-04, 9.9261886e-01, 8.2219703e-06, ..., 1.92755
10e-06,
        7.1251350e-03, 3.2941814e-05],
       ...,
       [1.0717752e-06, 7.5279139e-08, 4.6543805e-06, ..., 1.41682
93e-06,
        3.6145357e-07, 2.2929487e-07],
       [8.8150543e-04, 9.9872953e-01, 1.4206319e-04, ..., 9.55433
55e-06,
        5.1114512e-05, 4.8709891e-05],
       [2.7192047e-07, 1.7739443e-06, 4.2710103e-06, ..., 9.99820
53e-01,
        1.7729168e-07, 1.5285700e-07]], dtype=float32)
```

In [214]:



```
1 import numpy as np
2 test_predict_result = np.array([np.argmax(test_predict[i]) for i in range(1
3 test_predict_result
```

Out[214]:

```
array([3, 8, 1, ..., 5, 1, 7])
```

In [215]:



```
1 test_labels = np.array([np.argmax(cifar.test.labels[i])
2                       for i in range(len(cifar.test.labels))])
3 test_labels
```


Out[215]:

```
array([3, 8, 8, ..., 5, 1, 7])
```

In [217]:



```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 %matplotlib inline
4
5 from sklearn.metrics import confusion_matrix
6
7 mat = confusion_matrix(test_predict_result, test_labels)
8
9 plt.figure(figsize=(10,8))
10 sns.heatmap(mat, square=False, annot=True, fmt='d', cbar=True)
11 plt.xlim((0, 10))
12 plt.ylim((10, 0))
13 plt.xlabel('predicted value')
14 plt.ylabel('true value')
15 plt.show()
16
17 print(mat)
```



```
[[853  6 18 11  7  4  3  3 37  7]
```

```
[ 7 953 0 1 2 0 0 0 14 47]
[ 57 3 871 54 50 28 24 17 10 3]
[ 8 1 14 666 20 68 11 12 6 3]
[ 7 0 30 39 863 23 13 33 3 0]
[ 4 1 27 137 17 845 8 17 1 1]
[ 7 4 27 59 21 12 929 1 4 1]
[ 8 0 9 16 19 19 5 912 2 3]
[ 28 3 2 4 1 0 4 1 907 9]
[ 21 29 2 13 0 1 3 4 16 926]]
```