

## < Deep Learning - PART2 TF2 CNNs >

### Ch 5. CNNs Workshop 5 - Transfer learning with TF Hub ¶

2021/10/01

[ Reference ] :

1. TensorFlow Core - Tutorials: **Transfer learning with TensorFlow Hub**  
[https://www.tensorflow.org/tutorials/images/transfer\\_learning\\_with\\_hub?hl=zh\\_tw](https://www.tensorflow.org/tutorials/images/transfer_learning_with_hub?hl=zh_tw)  
([https://www.tensorflow.org/tutorials/images/transfer\\_learning\\_with\\_hub?hl=zh\\_tw](https://www.tensorflow.org/tutorials/images/transfer_learning_with_hub?hl=zh_tw))
2. Keras Documentation - **Keras Applications** <https://keras.io/applications/>  
(<https://keras.io/applications/>)
3. Andrew Ng - Coursera : **Transfer learning** <https://zh-tw.coursera.org/lecture/machine-learning-projects/transfer-learning-WNPap>  
(<https://zh-tw.coursera.org/lecture/machine-learning-projects/transfer-learning-WNPap>)

[TensorFlow Hub \(https://tfhub.dev/\)](https://tfhub.dev/) is a repository of pre-trained TensorFlow models.

This tutorial demonstrates how to:

1. Use models from TensorFlow Hub with `tf.keras`
2. Use an image classification model from TensorFlow Hub
3. Do simple transfer learning to fine-tune a model for your own image classes

## Setup

On **Anaconda Prompt** , run :

```
pip install tensorflow_hub
```

In [ ]:



```
1 import numpy as np
2 import time
3
4 import PIL.Image as Image
5 import matplotlib.pyplot as plt
6
7 import tensorflow as tf
8 import tensorflow_hub as hub
9
10 import datetime
11
12 %load_ext tensorboard
```

## An ImageNet classifier

You'll start by using a classifier model pre-trained on the [ImageNet](https://en.wikipedia.org/wiki/ImageNet) (<https://en.wikipedia.org/wiki/ImageNet>) benchmark dataset—no initial training required!

### Download the classifier

Select a [MobileNetV2](https://arxiv.org/abs/1801.04381) (<https://arxiv.org/abs/1801.04381>) pre-trained model [from TensorFlow Hub](https://tfhub.dev/google/tf2-preview/mobilenet_v2/classification/2) ([https://tfhub.dev/google/tf2-preview/mobilenet\\_v2/classification/2](https://tfhub.dev/google/tf2-preview/mobilenet_v2/classification/2)) and wrap it as a Keras layer with `hub.KerasLayer` ([https://www.tensorflow.org/hub/api\\_docs/python/hub/KerasLayer](https://www.tensorflow.org/hub/api_docs/python/hub/KerasLayer)). Any [compatible image classifier model](https://tfhub.dev/s?q=tf2&module-type=image-classification/) (<https://tfhub.dev/s?q=tf2&module-type=image-classification/>) from TensorFlow Hub will work here, including the examples provided in the drop-down below.

In [4]:



```
1 mobilenet_v2 = "https://tfhub.dev/google/tf2-preview/mobilenet_v2/classifica
2 inception_v3 = "https://tfhub.dev/google/imagenet/inception_v3/classificati
3
4 classifier_model = mobilenet_v2 #@param ["mobilenet_v2", "inception_v3"] {t
```

In [5]:



```
1 IMAGE_SHAPE = (224, 224)
2
3 classifier = tf.keras.Sequential([
4     hub.KerasLayer(classifier_model, input_shape=IMAGE_SHAPE+(3,))
5 ])
```

WARNING:tensorflow:AutoGraph could not transform <bound method KerasLayer.call of <tensorflow\_hub.keras\_layer.KerasLayer object at 0x000002247099E760>> and will run it as-is.

Please report this to the TensorFlow team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH\_VERBOSITY=10`) and attach the full output.

Cause: module 'gast' has no attribute 'Index'

To silence this warning, decorate the function with @tf.autograph.experimental.do\_not\_convert

WARNING:tensorflow:AutoGraph could not transform <bound method KerasLayer.call of <tensorflow\_hub.keras\_layer.KerasLayer object at 0x000002247099E760>> and will run it as-is.

Please report this to the TensorFlow team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH\_VERBOSITY=10`) and attach the full output.

Cause: module 'gast' has no attribute 'Index'

To silence this warning, decorate the function with @tf.autograph.experimental.do\_not\_convert

WARNING: AutoGraph could not transform <bound method KerasLayer.call of <tensorflow\_hub.keras\_layer.KerasLayer object at 0x000002247099E760>> and will run it as-is.

Please report this to the TensorFlow team. When filing the bug, set the verbosity to 10 (on Linux, `export AUTOGRAPH\_VERBOSITY=10`) and attach the full output.

Cause: module 'gast' has no attribute 'Index'

To silence this warning, decorate the function with @tf.autograph.experimental.do\_not\_convert

## Run it on a single image

Download a single image to try the model on:

In [6]:

```
1 grace_hopper = tf.keras.utils.get_file('image.jpg', 'https://storage.googleapis.com/keras-datasets/image.jpg')
2 grace_hopper = Image.open(grace_hopper).resize(IMAGE_SHAPE)
3 grace_hopper
```

Out[6]:



In [7]:

```
1 grace_hopper = np.array(grace_hopper)/255.0
2 grace_hopper.shape
```

Out[7]:

(224, 224, 3)

Add a batch dimension (with `np.newaxis` ) and pass the image to the model:

In [8]:

```
1 result = classifier.predict(grace_hopper[np.newaxis, ...])
2 result.shape
```

Out[8]:

(1, 1001)

The result is a 1001-element vector of logits, rating the probability of each class for the image.

The top class ID can be found with `tf.math.argmax` :

In [9]:

```
1 predicted_class = tf.math.argmax(result[0], axis=-1)
2 predicted_class
```

Out[9]:

```
<tf.Tensor: shape=(), dtype=int64, numpy=653>
```

## Decode the predictions

Take the `predicted_class` ID (such as 653 ) and fetch the ImageNet dataset labels to decode the predictions:

In [10]:

```
1 labels_path = tf.keras.utils.get_file('ImageNetLabels.txt', 'https://storage
2 imagenet_labels = np.array(open(labels_path).read().splitlines())
```

In [11]:

```
1 plt.imshow(grace_hopper)
2 plt.axis('off')
3 predicted_class_name = imagenet_labels[predicted_class]
4 _ = plt.title("Prediction: " + predicted_class_name.title())
```



## Simple transfer learning

But what if you want to create a custom classifier using your own dataset that has classes that aren't included in the original ImageNet dataset (that the pre-trained model was trained on)?

To do that, you can:

1. Select a pre-trained model from TensorFlow Hub; and
2. Retrain the top (last) layer to recognize the classes from your custom dataset.

## Dataset

In this example, you will use the TensorFlow flowers dataset:

In [12]:



```
1 data_root = tf.keras.utils.get_file(  
2     'flower_photos',  
3     'https://storage.googleapis.com/download.tensorflow.org/example_images/fl  
4     untar=True)
```

First, load this data into the model using the image data off disk with

`tf.keras.preprocessing.image_dataset_from_directory`, which will generate a `tf.data.Dataset`:

In [13]:



```
1 batch_size = 32  
2 img_height = 224  
3 img_width = 224  
4  
5 train_ds = tf.keras.preprocessing.image_dataset_from_directory(  
6     str(data_root),  
7     validation_split=0.2,  
8     subset="training",  
9     seed=123,  
10    image_size=(img_height, img_width),  
11    batch_size=batch_size  
12 )  
13  
14 val_ds = tf.keras.preprocessing.image_dataset_from_directory(  
15     str(data_root),  
16     validation_split=0.2,  
17     subset="validation",  
18     seed=123,  
19     image_size=(img_height, img_width),  
20     batch_size=batch_size  
21 )
```

```
Found 3670 files belonging to 5 classes.  
Using 2936 files for training.  
Found 3670 files belonging to 5 classes.  
Using 734 files for validation.
```

The flowers dataset has five classes:

In [14]:



```
1 class_names = np.array(train_ds.class_names)
2 print(class_names)
```

```
['daisy' 'dandelion' 'roses' 'sunflowers' 'tulips']
```

Second, because TensorFlow Hub's convention for image models is to expect float inputs in the `[0, 1]` range, use the `tf.keras.layers.experimental.preprocessing.Rescaling` layer to achieve this.

Note: You could also include the

`tf.keras.layers.experimental.preprocessing.Rescaling` layer inside the model. Refer to the [Working with preprocessing layers](https://www.tensorflow.org/guide/keras/preprocessing_layers) ([https://www.tensorflow.org/guide/keras/preprocessing\\_layers](https://www.tensorflow.org/guide/keras/preprocessing_layers)) guide for a discussion of the tradeoffs.

In [15]:



```
1 normalization_layer = tf.keras.layers.experimental.preprocessing.Rescaling(
2 train_ds = train_ds.map(lambda x, y: (normalization_layer(x), y)) # Where x
3 val_ds = val_ds.map(lambda x, y: (normalization_layer(x), y)) # Where x=ima
```

Third, finish the input pipeline by using buffered prefetching with `Dataset.prefetch`, so you can yield the data from disk without I/O blocking issues.

These are some of the most important `tf.data` methods you should use when loading data. Interested readers can learn more about them, as well as how to cache data to disk and other techniques, in the [Better performance with the tf.data API](https://www.tensorflow.org/guide/data_performance#prefetching) ([https://www.tensorflow.org/guide/data\\_performance#prefetching](https://www.tensorflow.org/guide/data_performance#prefetching)) guide.

In [16]:



```
1 AUTOTUNE = tf.data.AUTOTUNE
2 train_ds = train_ds.cache().prefetch(buffer_size=AUTOTUNE)
3 val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

In [17]:



```
1 for image_batch, labels_batch in train_ds:
2     print(image_batch.shape)
3     print(labels_batch.shape)
4     break
```

```
(32, 224, 224, 3)
(32,)
```

## Run the classifier on a batch of images

Now, run the classifier on an image batch:

In [18]:



```
1 result_batch = classifier.predict(train_ds)
```

In [19]:



```
1 predicted_class_names = imagenet_labels[tf.math.argmax(result_batch, axis=-1)]
2 predicted_class_names
```

Out[19]:

```
array(['daisy', 'coral fungus', 'rapeseed', ..., 'daisy', 'daisy',
      'birdhouse'], dtype='<U30')
```

Check how these predictions line up with the images:

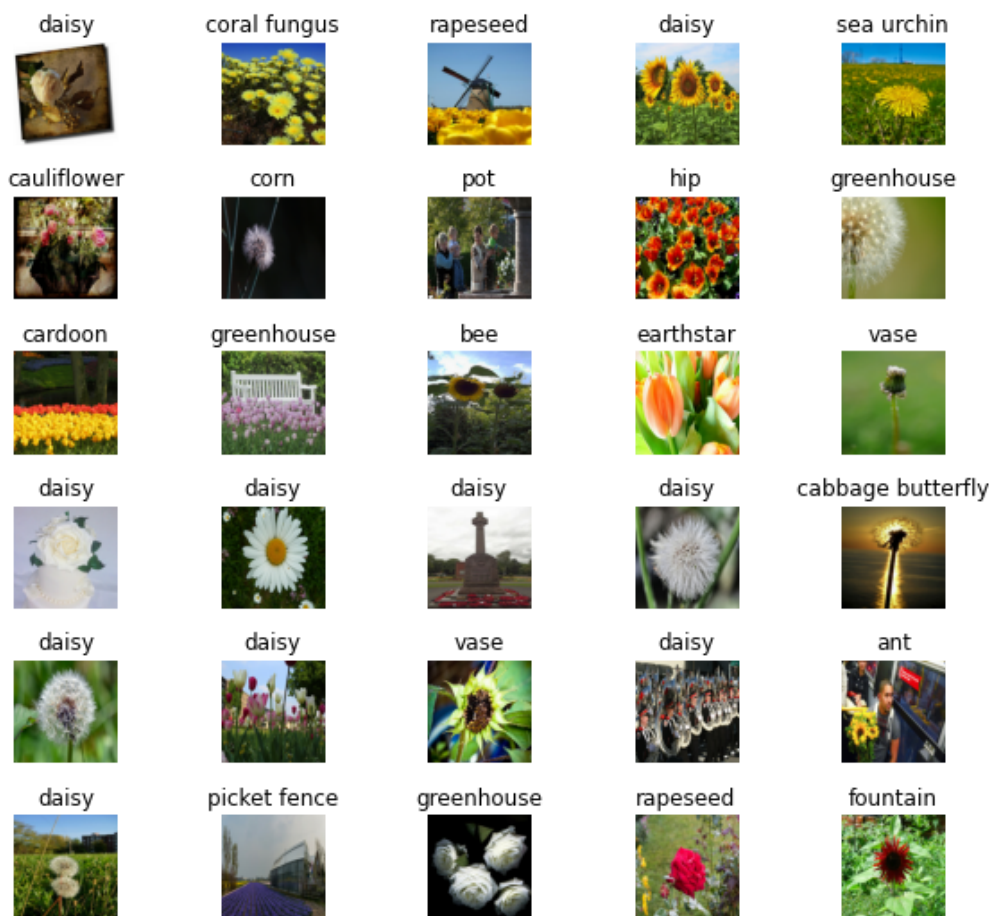


In [20]:



```
1 plt.figure(figsize=(10,9))
2 plt.subplots_adjust(hspace=0.5)
3 for n in range(30):
4     plt.subplot(6,5,n+1)
5     plt.imshow(image_batch[n])
6     plt.title(predicted_class_names[n])
7     plt.axis('off')
8 _ = plt.suptitle("ImageNet predictions")
```

ImageNet predictions



Note: all images are licensed CC-BY, creators are listed in the LICENSE.txt file.

The results are far from perfect, but reasonable considering that these are not the classes the model was trained for (except for "daisy").

## Download the headless model

TensorFlow Hub also distributes models without the top classification layer. These can be used to easily perform transfer learning.

Select a [MobileNetV2 \(https://arxiv.org/abs/1801.04381\)](https://arxiv.org/abs/1801.04381) pre-trained model from [TensorFlow Hub \(https://tfhub.dev/google/tf2-preview/mobilenet\\_v2/feature\\_vector/4\)](https://tfhub.dev/google/tf2-preview/mobilenet_v2/feature_vector/4). Any [compatible image feature vector model \(https://tfhub.dev/s?module-type=image-feature-vector&q=tf2\)](https://tfhub.dev/s?module-type=image-feature-vector&q=tf2) from TensorFlow Hub will work here, including the examples from the drop-down menu.

In [21]:

```
1 mobilenet_v2 = "https://tfhub.dev/google/tf2-preview/mobilenet_v2/feature_v
2 inception_v3 = "https://tfhub.dev/google/tf2-preview/inception_v3/feature_v
3
4 feature_extractor_model = mobilenet_v2 #@param ["mobilenet_v2", "inception_v3"]
```

Create the feature extractor by wrapping the pre-trained model as a Keras layer with `hub.KerasLayer` ([https://www.tensorflow.org/hub/api\\_docs/python/hub/KerasLayer](https://www.tensorflow.org/hub/api_docs/python/hub/KerasLayer)). Use the `trainable=False` argument to freeze the variables, so that the training only modifies the new classifier layer:

In [22]:

```
1 feature_extractor_layer = hub.KerasLayer(
2     feature_extractor_model,
3     input_shape=(224, 224, 3),
4     trainable=False)
```

The feature extractor returns a 1280-long vector for each image (the image batch size remains at 32 in this example):

In [23]:

```
1 feature_batch = feature_extractor_layer(image_batch)
2 print(feature_batch.shape)
```

(32, 1280)

## Attach a classification head

To complete the model, wrap the feature extractor layer in a `tf.keras.Sequential` model and add a fully-connected layer for classification:

In [24]:



```
1 num_classes = len(class_names)
2
3 model = tf.keras.Sequential([
4     feature_extractor_layer,
5     tf.keras.layers.Dense(num_classes)
6 ])
7
8 model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
keras_layer_1 (KerasLayer)	(None, 1280)	2257984
dense (Dense)	(None, 5)	6405
=====		
Total params: 2,264,389		
Trainable params: 6,405		
Non-trainable params: 2,257,984		

In [25]:



```
1 predictions = model(image_batch)
```

In [26]:



```
1 predictions.shape
```

Out[26]:

TensorShape([32, 5])

## Train the model

Use `Model.compile` to configure the training process and add a `tf.keras.callbacks.TensorBoard` callback to create and store logs:

In [27]:



```
1 model.compile(  
2     optimizer=tf.keras.optimizers.Adam(),  
3     loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
4     metrics=['acc'])  
5  
6 log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")  
7 tensorboard_callback = tf.keras.callbacks.TensorBoard(  
8     log_dir=log_dir,  
9     histogram_freq=1) # Enable histogram computation for every epoch.
```

Now use the `Model.fit` method to train the model.

To keep this example short, you'll be training for just 10 epochs. To visualize the training progress in TensorBoard later, create and store logs as a [TensorBoard callback](https://www.tensorflow.org/tensorboard/get_started#using_tensorboard_with_keras_modelfit) ([https://www.tensorflow.org/tensorboard/get\\_started#using\\_tensorboard\\_with\\_keras\\_modelfit](https://www.tensorflow.org/tensorboard/get_started#using_tensorboard_with_keras_modelfit)).

In [28]:



```
1 NUM_EPOCHS = 10
2
3 history = model.fit(train_ds,
4                     validation_data=val_ds,
5                     epochs=NUM_EPOCHS,
6                     callbacks=tensorboard_callback)
```

Epoch 1/10

92/92 [=====] - 67s 695ms/step - loss: 1.1070 - acc: 0.5653 - val\_loss: 0.4672 - val\_acc: 0.8474

Epoch 2/10

92/92 [=====] - 71s 778ms/step - loss: 0.3976 - acc: 0.8619 - val\_loss: 0.3764 - val\_acc: 0.8706

Epoch 3/10

92/92 [=====] - 82s 896ms/step - loss: 0.3058 - acc: 0.9010 - val\_loss: 0.3406 - val\_acc: 0.8842

Epoch 4/10

92/92 [=====] - 96s 1s/step - loss: 0.2521 - acc: 0.9248 - val\_loss: 0.3221 - val\_acc: 0.8924

Epoch 5/10

92/92 [=====] - 104s 1s/step - loss: 0.2146 - acc: 0.9380 - val\_loss: 0.3108 - val\_acc: 0.8951

Epoch 6/10

92/92 [=====] - 111s 1s/step - loss: 0.1863 - acc: 0.9463 - val\_loss: 0.3032 - val\_acc: 0.8978

Epoch 7/10

92/92 [=====] - 105s 1s/step - loss: 0.1641 - acc: 0.9551 - val\_loss: 0.2979 - val\_acc: 0.9005

Epoch 8/10

92/92 [=====] - 103s 1s/step - loss: 0.1460 - acc: 0.9643 - val\_loss: 0.2939 - val\_acc: 0.9005

Epoch 9/10

92/92 [=====] - 101s 1s/step - loss: 0.1308 - acc: 0.9717 - val\_loss: 0.2909 - val\_acc: 0.9019

Epoch 10/10

92/92 [=====] - 105s 1s/step - loss: 0.1179 - acc: 0.9739 - val\_loss: 0.2885 - val\_acc: 0.9046

Start the TensorBoard to view how the metrics change with each epoch and to track other scalar values:

In [ ]:



```
1 # %tensorboard --logdir Logs/fit
```

**To run TensorBoard, run the following command on Anaconda (Powershell) Prompt :**

```
tensorboard --logdir= path/to/log-directory
```

- For instance, `tensorboard --logdir logs/fit`

Connecting to `http://localhost:6006`

## Check the predictions

Obtain the ordered list of class names from the model predictions:

In [30]:



```
1 predicted_batch = model.predict(image_batch)
2 predicted_id = tf.math.argmax(predicted_batch, axis=-1)
3 predicted_label_batch = class_names[predicted_id]
4 print(predicted_label_batch)
```

```
['roses' 'dandelion' 'tulips' 'sunflowers' 'dandelion' 'roses' 'dandelion'
 'roses' 'tulips' 'dandelion' 'tulips' 'tulips' 'sunflowers' 'tulips'
 'dandelion' 'roses' 'daisy' 'tulips' 'dandelion' 'dandelion' 'dandelion'
 'tulips' 'sunflowers' 'roses' 'sunflowers' 'dandelion' 'tulips'
 'roses'
 'roses' 'sunflowers' 'tulips' 'sunflowers']
```

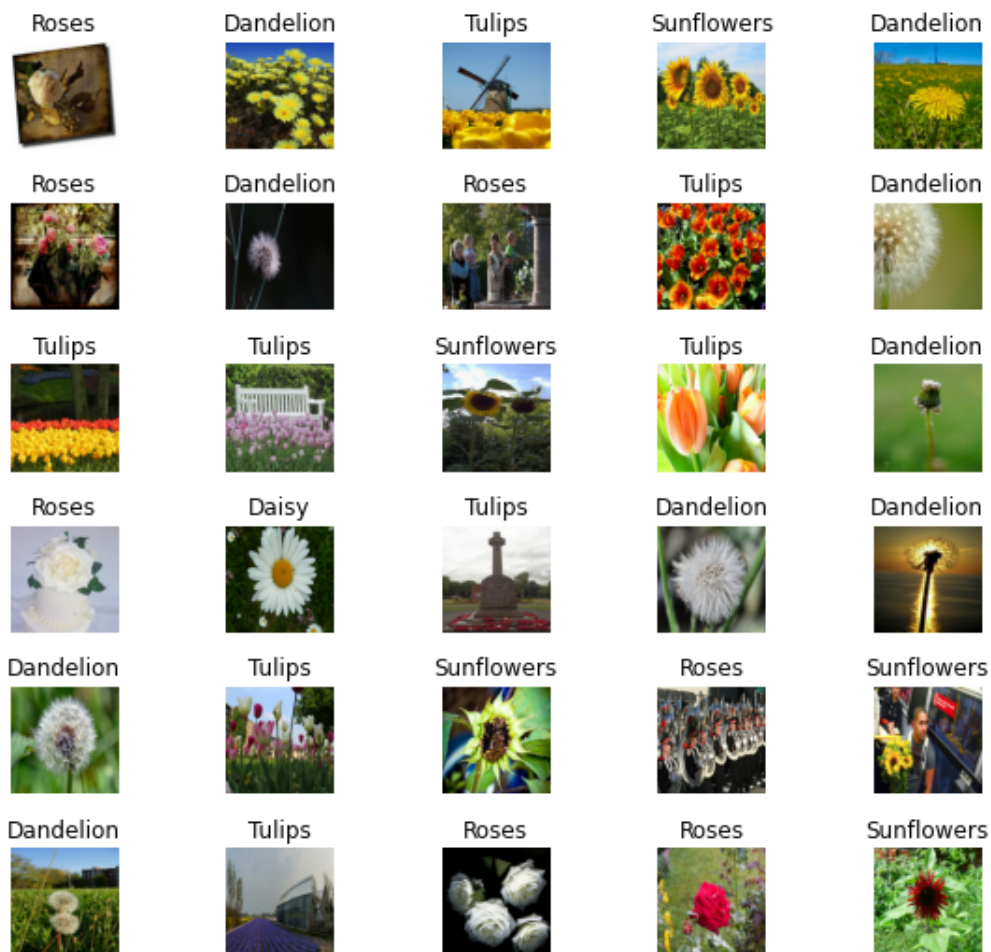
Plot the model predictions:

In [31]:



```
1 plt.figure(figsize=(10,9))
2 plt.subplots_adjust(hspace=0.5)
3
4 for n in range(30):
5     plt.subplot(6,5,n+1)
6     plt.imshow(image_batch[n])
7     plt.title(predicted_label_batch[n].title())
8     plt.axis('off')
9 _ = plt.suptitle("Model predictions")
```

Model predictions



## Export and reload your model

Now that you've trained the model, export it as a SavedModel for reusing it later.

In [32]:



```
1 t = time.time()
2
3 export_path = "/tmp/saved_models/{}".format(int(t))
4 model.save(export_path)
5
6 export_path
```

INFO:tensorflow:Assets written to: /tmp/saved\_models/1633297160\assets

INFO:tensorflow:Assets written to: /tmp/saved\_models/1633297160\assets

Out[32]:

'/tmp/saved\_models/1633297160'

Confirm that you can reload the SavedModel and that the model is able to output the same results:

In [33]:



```
1 reloaded = tf.keras.models.load_model(export_path)
```

In [34]:



```
1 result_batch = model.predict(image_batch)
2 reloaded_result_batch = reloaded.predict(image_batch)
```



In [35]:



```
1 abs(reloaded_result_batch - result_batch).max()
```

Out[35]:

0.0

In [36]:



```
1 reloaded_predicted_id = tf.math.argmax(reloaded_result_batch, axis=-1)
2 reloaded_predicted_label_batch = class_names[reloaded_predicted_id]
3 print(reloaded_predicted_label_batch)
```

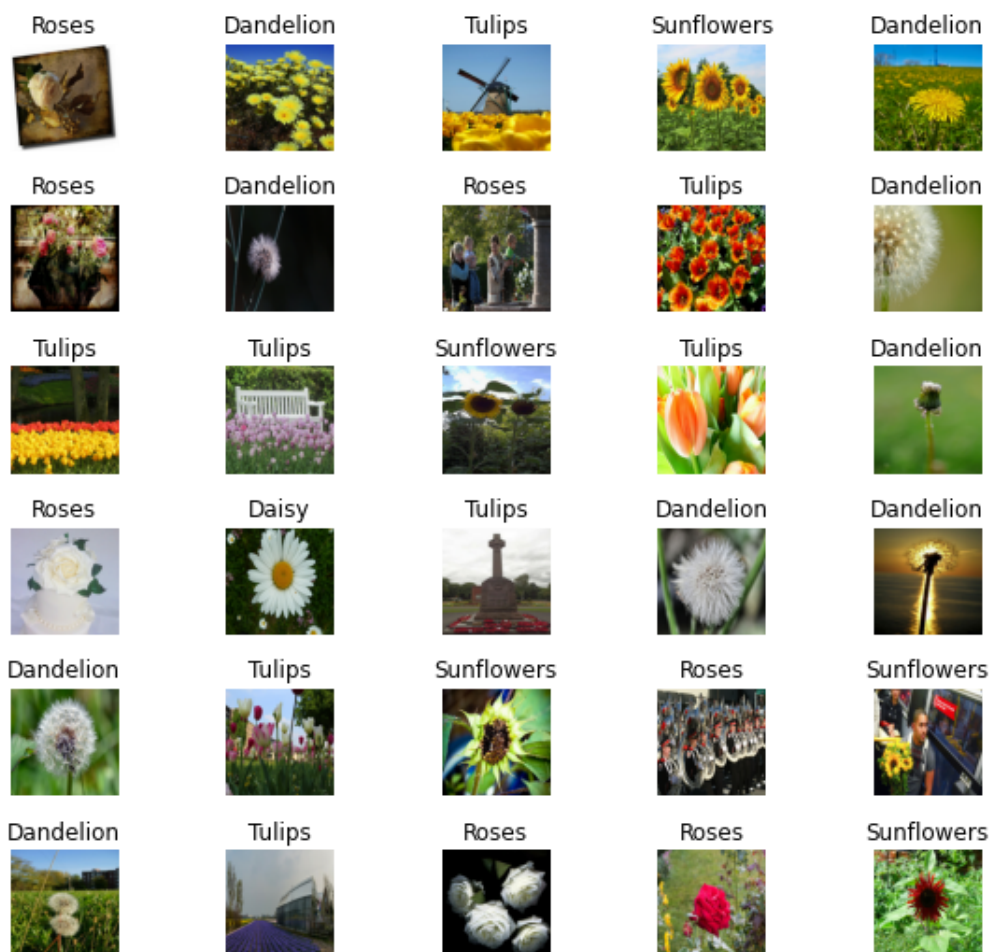
```
['roses' 'dandelion' 'tulips' 'sunflowers' 'dandelion' 'roses' 'dandelion'
 'roses' 'tulips' 'dandelion' 'tulips' 'tulips' 'sunflowers' 'tulips'
 'dandelion' 'roses' 'daisy' 'tulips' 'dandelion' 'dandelion' 'dandelion'
 'tulips' 'sunflowers' 'roses' 'sunflowers' 'dandelion' 'tulips'
 'roses'
 'roses' 'sunflowers' 'tulips' 'sunflowers']
```

In [37]:



```
1 plt.figure(figsize=(10,9))
2 plt.subplots_adjust(hspace=0.5)
3 for n in range(30):
4     plt.subplot(6,5,n+1)
5     plt.imshow(image_batch[n])
6     plt.title(reloaded_predicted_label_batch[n].title())
7     plt.axis('off')
8 _ = plt.suptitle("Model predictions")
```

Model predictions



## Next steps

You can use the SavedModel to load for inference or convert it to a [TensorFlow Lite](https://www.tensorflow.org/lite/convert/) model (for on-device machine learning) or a [TensorFlow.js](https://www.tensorflow.org/js/tutorials#convert_pretrained_models_to_tensorflowjs) model (for machine learning in JavaScript).

Discover [more tutorials](https://www.tensorflow.org/hub/tutorials) to learn how to use pre-trained models from TensorFlow Hub on image, text, audio, and video tasks.