

< Deep Learning - PART3 TF2 RNNs >

Ch 6. RNNs Workshop 3 - NLP - RNNs - IMDB : RNN & LSTM

2021/10/01

[Reference] :

- François Chollet, **Deep Learning with Python**, Chapter 6, Section 2, Manning, 2018.

[Code] : <https://github.com/fchollet/deep-learning-with-python-notebooks>
(<https://github.com/fchollet/deep-learning-with-python-notebooks>)

- Wikipedia, **Recurrent neural network**
https://en.wikipedia.org/wiki/Recurrent_neural_network
(https://en.wikipedia.org/wiki/Recurrent_neural_network)
- Wikipedia, **Long short-term memory**
https://en.wikipedia.org/wiki/Long_short-term_memory
(https://en.wikipedia.org/wiki/Long_short-term_memory)
- Andrej Karpathy blog, “**The Unreasonable Effectiveness of Recurrent Neural Networks**” <https://karpathy.github.io/2015/05/21/rnn-effectiveness/>
(<https://karpathy.github.io/2015/05/21/rnn-effectiveness/>)

[1. A first recurrent layer in tf.keras](#)

[2. A concrete LSTM example in tf.keras](#)

In [1]:

```
1 import tensorflow as tf
2 tf.__version__
```

Out[1]:

'2.4.1'

1. A first recurrent layer in tf.keras

The process we just naively implemented in Numpy corresponds to an actual Keras layer: the SimpleRNN layer:

In [2]:

```
1 from tensorflow.keras.layers import SimpleRNN
```

There is just one minor difference: SimpleRNN processes batches of sequences, like all other Keras layers, not just a single sequence like in our Numpy example. This means that it takes inputs of shape (batch_size, timesteps, input_features), rather than (timesteps, input_features).

Like all recurrent layers in Keras, SimpleRNN can be run in two different modes: it can return either the full sequences of successive outputs for each timestep (a 3D tensor of shape (batch_size, timesteps, output_features)), or it can return only the last output for each input sequence (a 2D tensor of shape (batch_size, output_features)). These two modes are controlled by the return_sequences constructor argument. Let's take a look at an example:

2021/09/01 CSCGroup#

In [4]:

```
1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Embedding, SimpleRNN
3
4 model = Sequential()
5 model.add(Embedding(10000, 32))
6 model.add(SimpleRNN(16))
7 model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
embedding_1 (Embedding)	(None, None, 32)	320000
=====		
simple_rnn_1 (SimpleRNN)	(None, 16)	784
=====		
Total params: 320,784		
Trainable params: 320,784		
Non-trainable params: 0		

In [4]:



```
1 model = Sequential()
2 model.add(Embedding(10000, 32))
3 model.add(SimpleRNN(32, return_sequences=True))
4 model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
embedding_1 (Embedding)	(None, None, 32)	320000

simple_rnn_1 (SimpleRNN)	(None, None, 32)	2080
=====		
Total params: 322,080		
Trainable params: 322,080		
Non-trainable params: 0		

It is sometimes useful to stack several recurrent layers one after the other in order to increase the representational power of a network. In such a setup, you have to get all intermediate layers to return full sequences:

In [5]:



```
1 model = Sequential()
2 model.add(Embedding(10000, 32))
3 model.add(SimpleRNN(32, return_sequences=True))
4 model.add(SimpleRNN(32, return_sequences=True))
5 model.add(SimpleRNN(32, return_sequences=True))
6 model.add(SimpleRNN(32)) # This last layer only returns the last outputs.
7 model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
embedding_2 (Embedding)	(None, None, 32)	320000

simple_rnn_2 (SimpleRNN)	(None, None, 32)	2080

simple_rnn_3 (SimpleRNN)	(None, None, 32)	2080

simple_rnn_4 (SimpleRNN)	(None, None, 32)	2080

simple_rnn_5 (SimpleRNN)	(None, 32)	2080
=====		
Total params: 328,320		
Trainable params: 328,320		
Non-trainable params: 0		

Now let's try to use such a model on the IMDB movie review classification problem. First, let's preprocess the data:

In [6]:

```
1 from tensorflow.keras.datasets import imdb
2 from tensorflow.keras.preprocessing import sequence
3
4 max_features = 10000 # number of words to consider as features
5 maxlen = 500 # cut texts after this number of words (among top max_features
6 batch_size = 32
7
8 print('Loading data...')
9 (input_train, y_train), (input_test, y_test) = imdb.load_data(num_words=max
10 print(len(input_train), 'train sequences')
11 print(len(input_test), 'test sequences')
12
13 print('Pad sequences (samples x time)')
14 input_train = sequence.pad_sequences(input_train, maxlen=maxlen)
15 input_test = sequence.pad_sequences(input_test, maxlen=maxlen)
16 print('input_train shape:', input_train.shape)
17 print('input_test shape:', input_test.shape)
```

Loading data...

```
<__array_function__ internals>:5: VisibleDeprecationWarning: Crea
ting an ndarray from ragged nested sequences (which is a list-or-
tuple of lists-or-tuples-or ndarrays with different lengths or sh
apes) is deprecated. If you meant to do this, you must specify 'd
type=object' when creating the ndarray
```

```
C:\Users\appcl\anaconda3\lib\site-packages\tensorflow\python\kera
s\datasets\imdb.py:159: VisibleDeprecationWarning: Creating an nd
array from ragged nested sequences (which is a list-or-tuple of l
ists-or-tuples-or ndarrays with different lengths or shapes) is d
eprecated. If you meant to do this, you must specify 'dtype=objec
t' when creating the ndarray
```

```
    x_train, y_train = np.array(xs[:idx]), np.array(labels[:idx])
```

```
C:\Users\appcl\anaconda3\lib\site-packages\tensorflow\python\kera
s\datasets\imdb.py:160: VisibleDeprecationWarning: Creating an nd
array from ragged nested sequences (which is a list-or-tuple of l
ists-or-tuples-or ndarrays with different lengths or shapes) is d
eprecated. If you meant to do this, you must specify 'dtype=objec
t' when creating the ndarray
```

```
    x_test, y_test = np.array(xs[idx:]), np.array(labels[idx:])
```

```
25000 train sequences
```

```
25000 test sequences
```

```
Pad sequences (samples x time)
```

```
input_train shape: (25000, 500)
```

```
input_test shape: (25000, 500)
```

Let's train a simple recurrent network using an Embedding layer and a SimpleRNN layer:

In [16]:



```
1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Embedding
3 from tensorflow.keras.layers import Dense
4
5 from tensorflow.keras import regularizers
6
7 model = Sequential()
8 model.add(Embedding(max_features, 8))
9 model.add(SimpleRNN(32, return_sequences=True,
10                     kernel_regularizer=regularizers.l2(0.001),
11                     recurrent_regularizer=regularizers.l2(0.001),
12                     bias_regularizer=None))
13 model.add(SimpleRNN(32, return_sequences=True,
14                     kernel_regularizer=regularizers.l2(0.001),
15                     recurrent_regularizer=regularizers.l2(0.001),
16                     bias_regularizer=None))
17 model.add(SimpleRNN(32, return_sequences=True,
18                     kernel_regularizer=regularizers.l2(0.001),
19                     recurrent_regularizer=regularizers.l2(0.001),
20                     bias_regularizer=None))model.add(SimpleRNN(16)) # This
21 model.add(Dense(1, activation='sigmoid'))
22
23 model.summary()
24
25 model.compile(optimizer='rmsprop',
26               loss='binary_crossentropy',
27               metrics=['acc'])
28 history = model.fit(input_train, y_train,
29                     epochs=10,
30                     batch_size=128,
31                     validation_split=0.2)
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
embedding_5 (Embedding)	(None, None, 8)	80000
simple_rnn_8 (SimpleRNN)	(None, None, 32)	1312
simple_rnn_9 (SimpleRNN)	(None, 16)	784
dense_1 (Dense)	(None, 1)	17

Total params: 82,113
Trainable params: 82,113

Non-trainable params: 0

```
Epoch 1/10
157/157 [=====] - 22s 129ms/step - loss: 0.7441 - acc: 0.5133 - val_loss: 0.7747 - val_acc: 0.5210
Epoch 2/10
157/157 [=====] - 22s 140ms/step - loss: 0.5775 - acc: 0.7333 - val_loss: 0.4326 - val_acc: 0.8302
Epoch 3/10
157/157 [=====] - 22s 141ms/step - loss: 0.4000 - acc: 0.8503 - val_loss: 0.4062 - val_acc: 0.8506
Epoch 4/10
157/157 [=====] - 22s 141ms/step - loss: 0.3251 - acc: 0.8911 - val_loss: 0.5188 - val_acc: 0.7922
Epoch 5/10
157/157 [=====] - 22s 143ms/step - loss: 0.3011 - acc: 0.9025 - val_loss: 0.4674 - val_acc: 0.8072
Epoch 6/10
157/157 [=====] - 22s 140ms/step - loss: 0.2335 - acc: 0.9321 - val_loss: 0.3809 - val_acc: 0.8766
Epoch 7/10
157/157 [=====] - 22s 141ms/step - loss: 0.2125 - acc: 0.9383 - val_loss: 0.4532 - val_acc: 0.8242
Epoch 8/10
157/157 [=====] - 21s 135ms/step - loss: 0.2670 - acc: 0.9129 - val_loss: 0.4915 - val_acc: 0.8214
Epoch 9/10
157/157 [=====] - 21s 133ms/step - loss: 0.1783 - acc: 0.9501 - val_loss: 0.5961 - val_acc: 0.8456
Epoch 10/10
157/157 [=====] - 21s 136ms/step - loss: 0.1487 - acc: 0.9635 - val_loss: 0.5767 - val_acc: 0.8160
```

Let's display the training and validation loss and accuracy:

In [17]:




```
1 print(history.history.keys())
```

```
dict_keys(['loss', 'acc', 'val_loss', 'val_acc'])
```

In [18]:



```
1 import matplotlib.pyplot as plt
2 %matplotlib inline
3
4 acc = history.history['acc']
5 val_acc = history.history['val_acc']
6 loss = history.history['loss']
7 val_loss = history.history['val_loss']
8
9 epochs = range(len(acc))
10
11 plt.plot(epochs, acc, 'bo', label='Training acc')
12 plt.plot(epochs, val_acc, 'b', label='Validation acc')
13 plt.title('Training and validation accuracy')
14 plt.legend()
15
16 plt.figure()
17
18 plt.plot(epochs, loss, 'bo', label='Training loss')
19 plt.plot(epochs, val_loss, 'b', label='Validation loss')
20 plt.title('Training and validation loss')
21 plt.legend()
22
23 plt.show()
```





As a reminder, in chapter 3, our very first naive approach to this very dataset got us to 88% test accuracy. Unfortunately, our small recurrent network doesn't perform very well at all compared to this baseline (only up to 85% validation accuracy). Part of the problem is that our inputs only consider the first 500 words rather the full sequences -- hence our RNN has access to less information than our earlier baseline model. The remainder of the problem is simply that `SimpleRNN` isn't very good at processing long sequences, like text. Other types of recurrent layers perform much better. Let's take a look at some more advanced layers.

2. A concrete LSTM example in `tf.keras`

Now let's switch to more practical concerns: we will set up a model using a LSTM layer and train it on the IMDB data. Here's the network, similar to the one with `SimpleRNN` that we just presented. We only specify the output dimensionality of the LSTM layer, and leave every other argument (there are lots) to the Keras defaults. Keras has good defaults, and things will almost always "just work" without you having to spend time tuning parameters by hand.

In [22]:



```
1 from tensorflow.keras.layers import LSTM
2 from tensorflow.keras.models import Sequential
3 from tensorflow.keras.layers import Embedding, Dense
4
5 max_features = 10000
6 model = Sequential()
7 model.add(Embedding(max_features, 32))
8 model.add(LSTM(64, return_sequences=True,
9               kernel_regularizer=regularizers.l2(0.001),
10              recurrent_regularizer=regularizers.l2(0.001),
11              bias_regularizer=None))
12 model.add(LSTM(32))
13 model.add(Dense(1, activation='sigmoid'))
14 model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====		
embedding_7 (Embedding)	(None, None, 32)	320000

lstm_1 (LSTM)	(None, None, 64)	24832

lstm_2 (LSTM)	(None, 32)	12416

dense_3 (Dense)	(None, 1)	33
=====		
Total params: 357,281		
Trainable params: 357,281		
Non-trainable params: 0		

In [23]:



```
1 model.compile(optimizer='rmsprop',
2               loss='binary_crossentropy',
3               metrics=['acc'])
4 history = model.fit(input_train, y_train,
5                     epochs=10,
6                     batch_size=128,
7                     validation_split=0.2)
```

Epoch 1/10

157/157 [=====] - 74s 449ms/step - loss: 0.6668 - acc: 0.6503 - val_loss: 0.4367 - val_acc: 0.8114

Epoch 2/10

157/157 [=====] - 81s 518ms/step - loss: 0.3627 - acc: 0.8628 - val_loss: 0.3903 - val_acc: 0.8480

Epoch 3/10

157/157 [=====] - 81s 518ms/step - loss: 0.3013 - acc: 0.8960 - val_loss: 0.3631 - val_acc: 0.8634

Epoch 4/10

157/157 [=====] - 80s 513ms/step - loss: 0.2443 - acc: 0.9160 - val_loss: 0.4709 - val_acc: 0.8348

Epoch 5/10

157/157 [=====] - 85s 541ms/step - loss: 0.2199 - acc: 0.9242 - val_loss: 0.3678 - val_acc: 0.8800

Epoch 6/10

157/157 [=====] - 85s 539ms/step - loss: 0.1949 - acc: 0.9362 - val_loss: 0.4297 - val_acc: 0.8722

Epoch 7/10

157/157 [=====] - 90s 576ms/step - loss: 0.1954 - acc: 0.9383 - val_loss: 0.4261 - val_acc: 0.8760

Epoch 8/10

157/157 [=====] - 88s 564ms/step - loss: 0.1697 - acc: 0.9465 - val_loss: 0.3654 - val_acc: 0.8800

Epoch 9/10

157/157 [=====] - 94s 602ms/step - loss: 0.1497 - acc: 0.9539 - val_loss: 0.3832 - val_acc: 0.8596

Epoch 10/10

157/157 [=====] - 98s 627ms/step - loss: 0.1314 - acc: 0.9601 - val_loss: 0.4139 - val_acc: 0.8682

In [24]:



```
1 acc = history.history['acc']
2 val_acc = history.history['val_acc']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']
5
6 epochs = range(len(acc))
7
8 plt.plot(epochs, acc, 'bo', label='Training acc')
9 plt.plot(epochs, val_acc, 'b', label='Validation acc')
10 plt.title('Training and validation accuracy')
11 plt.legend()
12
13 plt.figure()
14
15 plt.plot(epochs, loss, 'bo', label='Training loss')
16 plt.plot(epochs, val_loss, 'b', label='Validation loss')
17 plt.title('Training and validation loss')
18 plt.legend()
19
20 plt.show()
```

