

< Deep Learning - PART3 TF2 RNNs >

Ch 6. RNNs Workshop 4 - NLP - RNNs - IMDB : Bidirectional LSTM & GRU

2021/10/01

[Reference] :

1. Keras Documentation, **Trains a Bidirectional LSTM on the IMDB sentiment classification task**
https://keras.io/zh/examples/imdb_bidirectional_lstm/
(https://keras.io/zh/examples/imdb_bidirectional_lstm/).
2. François Chollet, **Deep Learning with Python**, Chapter 6, Section 2, Manning, 2018. [Code] : <https://github.com/fchollet/deep-learning-with-python-notebooks> (<https://github.com/fchollet/deep-learning-with-python-notebooks>).
3. Wikipedia, **Recurrent neural network**
https://en.wikipedia.org/wiki/Recurrent_neural_network
(https://en.wikipedia.org/wiki/Recurrent_neural_network).
 - Wikipedia, **Long short-term memory**
https://en.wikipedia.org/wiki/Long_short-term_memory
(https://en.wikipedia.org/wiki/Long_short-term_memory).
4. Andrej Karpathy blog, "**The Unreasonable Effectiveness of Recurrent Neural Networks**" <https://karpathy.github.io/2015/05/21/rnn-effectiveness/>
(<https://karpathy.github.io/2015/05/21/rnn-effectiveness/>).
5. Wikipedia, **Bidirectional recurrent neural networks**
https://en.wikipedia.org/wiki/Bidirectional_recurrent_neural_networks
(https://en.wikipedia.org/wiki/Bidirectional_recurrent_neural_networks).
6. Wikipedia, **Gated recurrent unit**
https://en.wikipedia.org/wiki/Gated_recurrent_unit
(https://en.wikipedia.org/wiki/Gated_recurrent_unit).
7. Tom Hope, Yehezkel S. Resheff, and Itay Lieder, "**Learning TensorFlow : A Guide to Building Deep Learning Systems**", Chapter 6, O'Reilly, 2017. - [Code] : https://github.com/giser-yugang/Learning_TensorFlow
(https://github.com/giser-yugang/Learning_TensorFlow).

[1. IMDB with Bidirectional LSTM](#)

[2. IMDB with Bidirectional GRU](#)

Bidirectional RNN layers (LSTM & GRU)

- Bidirectional RNN layers consist of, in their basic form, is two ordinary RNN layers:
 - one layer that reads the sequence from left to right, and
 - another that reads from right to left.
 - Each yields a hidden representation, the left-to-right vector, and the right-to-left vector.
 - These are then concatenated into one vector.
 - The major advantage of this representation is its ability to capture the context of words from both directions, which enables richer understanding of natural language and the underlying semantics in text.
 - In practice, in complex tasks, it often leads to improved accuracy.

In [1]:



```
1 import tensorflow as tf
2 tf.__version__
```

Out[1]:

'2.0.0'

1. IMDB with Bidirectional LSTM

In [2]:



```
1 from __future__ import print_function
2 import numpy as np
3
4 from tensorflow.keras.preprocessing import sequence
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.layers import Dense, Dropout, Embedding, LSTM, Bidire
7 from tensorflow.keras.datasets import imdb
```

In [3]:

```
1 max_features = 20000
2 # cut texts after this number of words
3 # (among top max_features most common words)
4 maxlen = 100
5 batch_size = 256
```

Loading IMDB dataset

In [4]:

```
1 print('Loading data...')
2 (x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)
3 print(len(x_train), 'train sequences')
4 print(len(x_test), 'test sequences')
```

Loading data...

25000 train sequences

25000 test sequences

In [5]:

```
1 print('Pad sequences (samples x time)')
2 x_train = sequence.pad_sequences(x_train, maxlen=maxlen)
3 x_test = sequence.pad_sequences(x_test, maxlen=maxlen)
4 print('x_train shape:', x_train.shape)
5 print('x_test shape:', x_test.shape)
```

Pad sequences (samples x time)

x_train shape: (25000, 100)

x_test shape: (25000, 100)

In [6]:

```
1 y_train = np.array(y_train)
2 y_test = np.array(y_test)
```

Bidirectional LSTM

In [7]:

```
1 model = Sequential()
2 model.add(Embedding(max_features, 128, input_length=maxlen))
3 model.add(Bidirectional(LSTM(64)))
4 model.add(Dropout(0.5))
5 model.add(Dense(1, activation='sigmoid'))
```

In [8]:



```
1 # try using different optimizers and different optimizer configs
2 model.compile('adam', 'binary_crossentropy', metrics=['accuracy'])
```

In [9]:



```
1 print('Train...')
2 history = model.fit(x_train, y_train,
3                     batch_size=batch_size,
4                     epochs=10,
5                     validation_data=[x_test, y_test])
```

Train...

Train on 25000 samples, validate on 25000 samples

Epoch 1/10

25000/25000 [=====] - 53s 2ms/sample - loss: 0.4990 - accuracy: 0.7445 - val_loss: 0.3436 - val_accuracy: 0.8474

Epoch 2/10

25000/25000 [=====] - 37s 1ms/sample - loss: 0.2508 - accuracy: 0.9011 - val_loss: 0.3408 - val_accuracy: 0.8496

Epoch 3/10

25000/25000 [=====] - 37s 1ms/sample - loss: 0.1700 - accuracy: 0.9391 - val_loss: 0.3971 - val_accuracy: 0.8400

Epoch 4/10

25000/25000 [=====] - 37s 1ms/sample - loss: 0.1127 - accuracy: 0.9624 - val_loss: 0.4693 - val_accuracy: 0.8313

Epoch 5/10

25000/25000 [=====] - 37s 1ms/sample - loss: 0.0789 - accuracy: 0.9750 - val_loss: 0.6139 - val_accuracy: 0.8199

Epoch 6/10

25000/25000 [=====] - 37s 1ms/sample - loss: 0.0587 - accuracy: 0.9806 - val_loss: 0.6681 - val_accuracy: 0.8109

Epoch 7/10

25000/25000 [=====] - 37s 1ms/sample - loss: 0.0428 - accuracy: 0.9867 - val_loss: 0.7591 - val_accuracy: 0.8184

Epoch 8/10

25000/25000 [=====] - 37s 1ms/sample - loss: 0.0329 - accuracy: 0.9898 - val_loss: 0.8509 - val_accuracy: 0.8206

Epoch 9/10

25000/25000 [=====] - 37s 1ms/sample - loss: 0.0256 - accuracy: 0.9928 - val_loss: 0.9110 - val_accuracy: 0.8166

Epoch 10/10

25000/25000 [=====] - 37s 1ms/sample - loss: 0.0262 - accuracy: 0.9924 - val_loss: 0.7756 - val_accuracy: 0.8181

Evaluation

In [10]:



```
1 model.evaluate(x_test, y_test, verbose=2)
```

25000/1 - 21s - loss: 1.0246 - accuracy: 0.8181

Out[10]:

[0.7756469815397262, 0.81812]

In [11]:



```
1 history_dict = history.history  
2 history_dict.keys()
```


Out[11]:


dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

In [12]:



```
1 import matplotlib.pyplot as plt
2 %matplotlib inline
3
4 acc = history.history['accuracy']
5 val_acc = history.history['val_accuracy']
6 loss = history.history['loss']
7 val_loss = history.history['val_loss']
8
9 epochs = range(1, len(acc) + 1)
10
11 plt.plot(epochs, acc, 'bo', label='Training acc')
12 plt.plot(epochs, val_acc, 'b', label='Validation acc')
13 plt.title('Training and validation accuracy')
14 plt.legend()
15
16 plt.figure()
17
18 plt.plot(epochs, loss, 'bo', label='Training loss')
19 plt.plot(epochs, val_loss, 'b', label='Validation loss')
20 plt.title('Training and validation loss')
21 plt.legend()
22
23 plt.show()
```





2. IMDB with Bidirectional GRU (Gated Recurrent Unit)

- Gated recurrent unit (GRU) cells are a simplification of sorts of LSTM cells.
- They also have a memory mechanism, but with considerably fewer parameters than LSTM.
- They are often used when there is less available data, and are faster to compute.

Bidirectional GRU

In [13]:

```
1 from tensorflow.keras.layers import Dense, Dropout, Embedding, GRU, Bidirectional
2
3 gru_model = Sequential()
4 gru_model.add(Embedding(max_features, 128, input_length=maxlen))
5 gru_model.add(Bidirectional(GRU(64)))
6 gru_model.add(Dropout(0.5))
7 gru_model.add(Dense(1, activation='sigmoid'))
```

In [14]:

```
1 # try using different optimizers and different optimizer configs
2 gru_model.compile('adam', 'binary_crossentropy', metrics=['accuracy'])
```


In [15]:



```
1 print('Bidirectional GRU...')
2 gru_history = gru_model.fit(x_train, y_train,
3                             batch_size=batch_size,
4                             epochs=10,
5                             validation_data=[x_test, y_test])
```

Bidirectional GRU...

Train on 25000 samples, validate on 25000 samples

Epoch 1/10

25000/25000 [=====] - 36s 1ms/sample - loss: 0.5324 - accuracy: 0.7148 - val_loss: 0.3576 - val_accuracy: 0.8449

Epoch 2/10

25000/25000 [=====] - 33s 1ms/sample - loss: 0.2672 - accuracy: 0.8936 - val_loss: 0.3476 - val_accuracy: 0.8499

Epoch 3/10

25000/25000 [=====] - 34s 1ms/sample - loss: 0.1676 - accuracy: 0.9406 - val_loss: 0.4293 - val_accuracy: 0.8404

Epoch 4/10

25000/25000 [=====] - 33s 1ms/sample - loss: 0.1075 - accuracy: 0.9643 - val_loss: 0.4934 - val_accuracy: 0.8329

Epoch 5/10

25000/25000 [=====] - 33s 1ms/sample - loss: 0.0619 - accuracy: 0.9796 - val_loss: 0.6604 - val_accuracy: 0.8286

Epoch 6/10

25000/25000 [=====] - 34s 1ms/sample - loss: 0.0422 - accuracy: 0.9862 - val_loss: 0.7159 - val_accuracy: 0.8259

Epoch 7/10

25000/25000 [=====] - 33s 1ms/sample - loss: 0.0262 - accuracy: 0.9922 - val_loss: 0.8887 - val_accuracy: 0.8210

Epoch 8/10

25000/25000 [=====] - 33s 1ms/sample - loss: 0.0169 - accuracy: 0.9952 - val_loss: 0.9565 - val_accuracy: 0.8188

Epoch 9/10

25000/25000 [=====] - 33s 1ms/sample - loss: 0.0099 - accuracy: 0.9973 - val_loss: 0.9881 - val_accuracy: 0.8176

Epoch 10/10

25000/25000 [=====] - 33s 1ms/sample - loss: 0.0060 - accuracy: 0.9987 - val_loss: 1.0629 - val_accuracy: 0.8174

In [16]:



```
1 gru_model.evaluate(x_test, y_test, verbose=2)
```

25000/1 - 18s - loss: 1.2307 - accuracy: 0.8174


Out[16]:


```
[1.062934914600849, 0.81744]
```

In [17]:



```
1 import matplotlib.pyplot as plt
2 %matplotlib inline
3
4 acc = gru_history.history['accuracy']
5 val_acc = gru_history.history['val_accuracy']
6 loss = gru_history.history['loss']
7 val_loss = gru_history.history['val_loss']
8
9 epochs = range(1, len(acc) + 1)
10
11 plt.plot(epochs, acc, 'bo', label='Training acc')
12 plt.plot(epochs, val_acc, 'b', label='Validation acc')
13 plt.title('Training and validation accuracy')
14 plt.legend()
15
16 plt.figure()
17
18 plt.plot(epochs, loss, 'bo', label='Training loss')
19 plt.plot(epochs, val_loss, 'b', label='Validation loss')
20 plt.title('Training and validation loss')
21 plt.legend()
22
23 plt.show()
```





Q1: What are the differences between Bidir. LSTM & Bidir. GRU models?

Q2: How to improve performances of both Bidir. LSTM & Bidir. GRU models?

In []:



1